



# EE3662 DSP Lab

## Unit II – Visual Signal

Instructor: 黃朝宗、孫民

TA: 汪立偉、丁友鈞、鄭安佑、鄭敬儒

National Tsing Hua University  
Department of Electrical Engineering



# Syllabus

- 1<sup>st</sup> Week: 3-hr lecture + take-home lab (Lab 5)
- 2<sup>nd</sup> Week: 3-hr lab on hybrid image (Lab 6)
- 3<sup>rd</sup> Week: 3-hr lab on corner detection (Lab 7)
- 4<sup>th</sup> Week: 3-hr lab on seam carving (Lab 8)

\* MATLAB is used for all labs



# Lecture Outline

- |                          | Lab Map                          |
|--------------------------|----------------------------------|
| • Image and Video Basics | Pixel array manipulation (Lab 5) |
| • Image Filtering        | Image filtering (Lab 6)          |
| • Interest Point         | Image feature (Lab 7)            |
| • Seam Carving           | Image Manipulation (Lab 8)       |



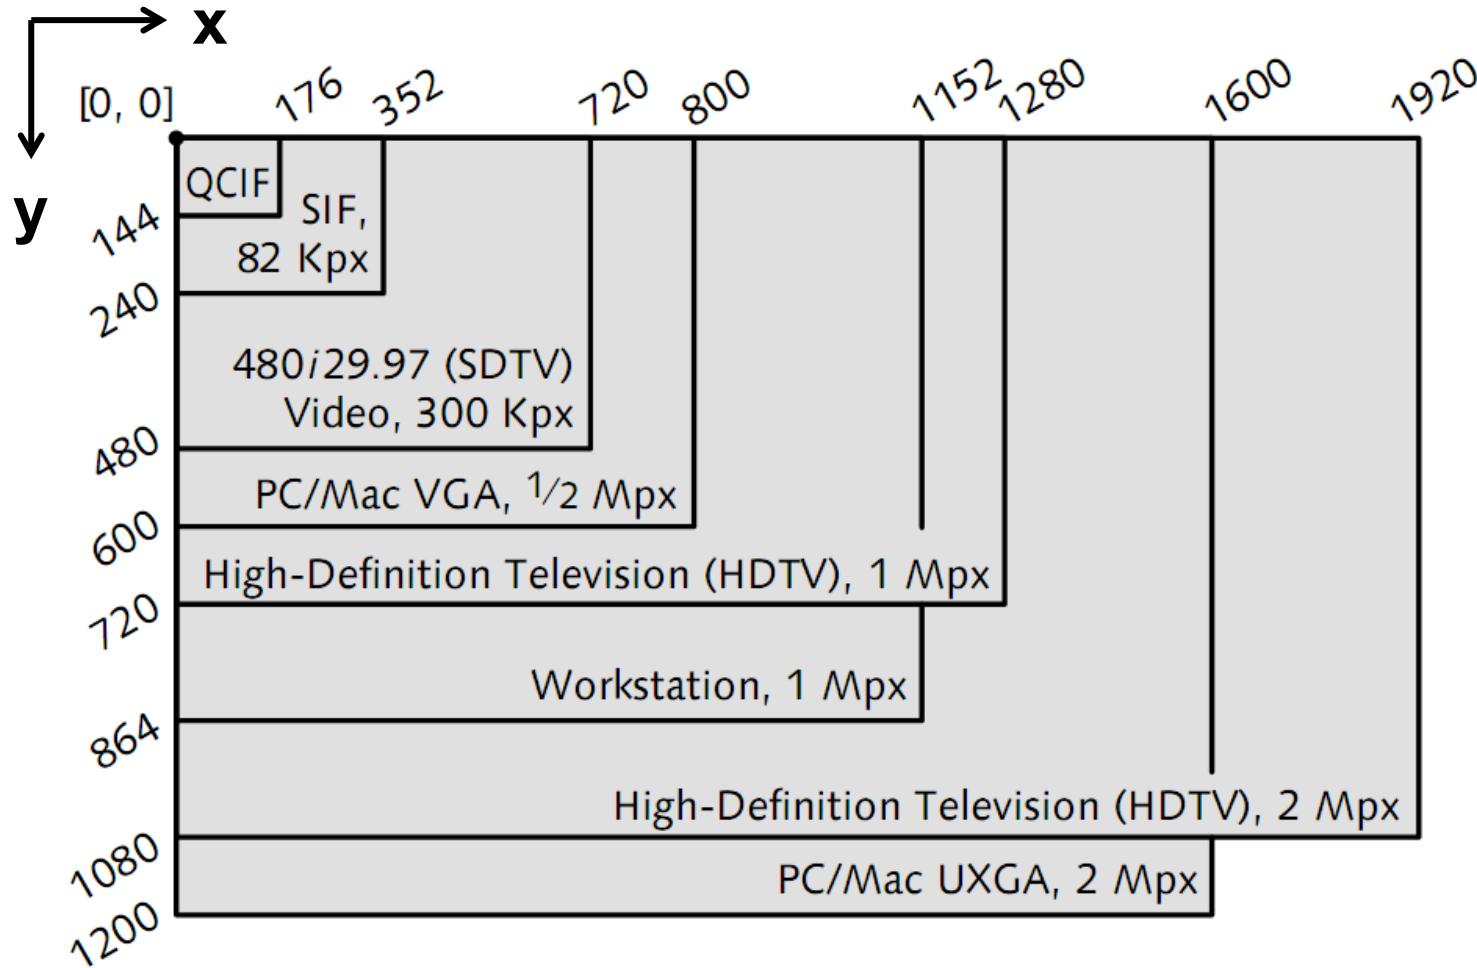
# Basics Representation

- Image
  - Pixel Array (spatial resolution)
  - Pixel Value (luminance and color)
- Video
  - Frame Rate (temporal resolution)



# Pixel Array

- Picture Element, Pel or Pixel, in (x,y)



Ref: *Digital Video and HDTV*

# Pixel Value

Grayscale (灰階) image



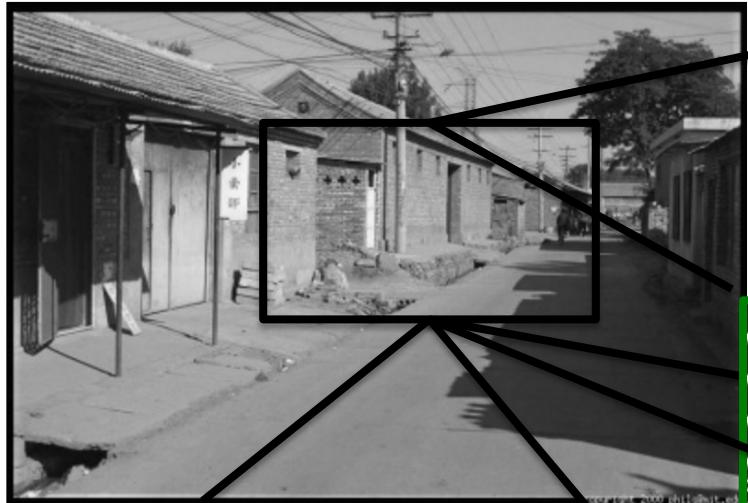
```
,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,  
0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,  
,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,  
,255,191,128,128,64,0,0,0,0,0,0,  
5,255,255,255,255,255,191,64,0,0,0  
55,255,255,255,255,255,255,255,255,  
,0,128,191,255,255,255,191,0,0,0,0  
,255,255,128,0,0,0,0,0,0,0,0,0,0,0,  
,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0
```

0,1,2,...255 uint8 values



# RGB color space

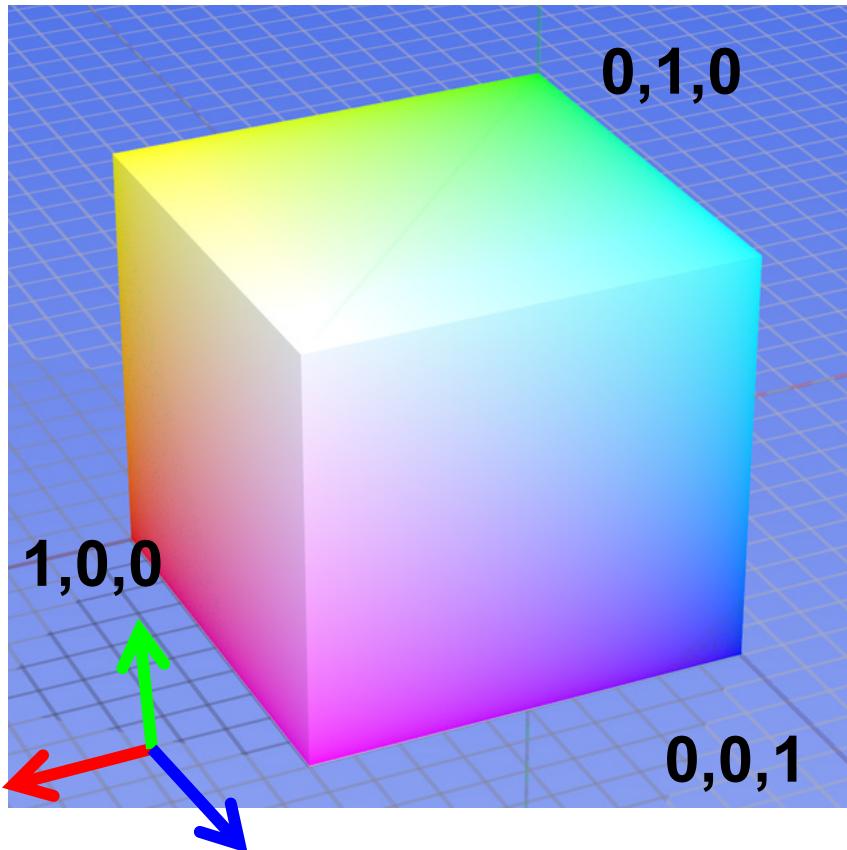
## Color image



# Color spaces: RGB



Default color space



Some drawbacks

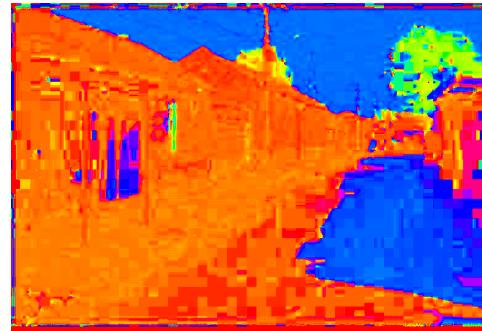
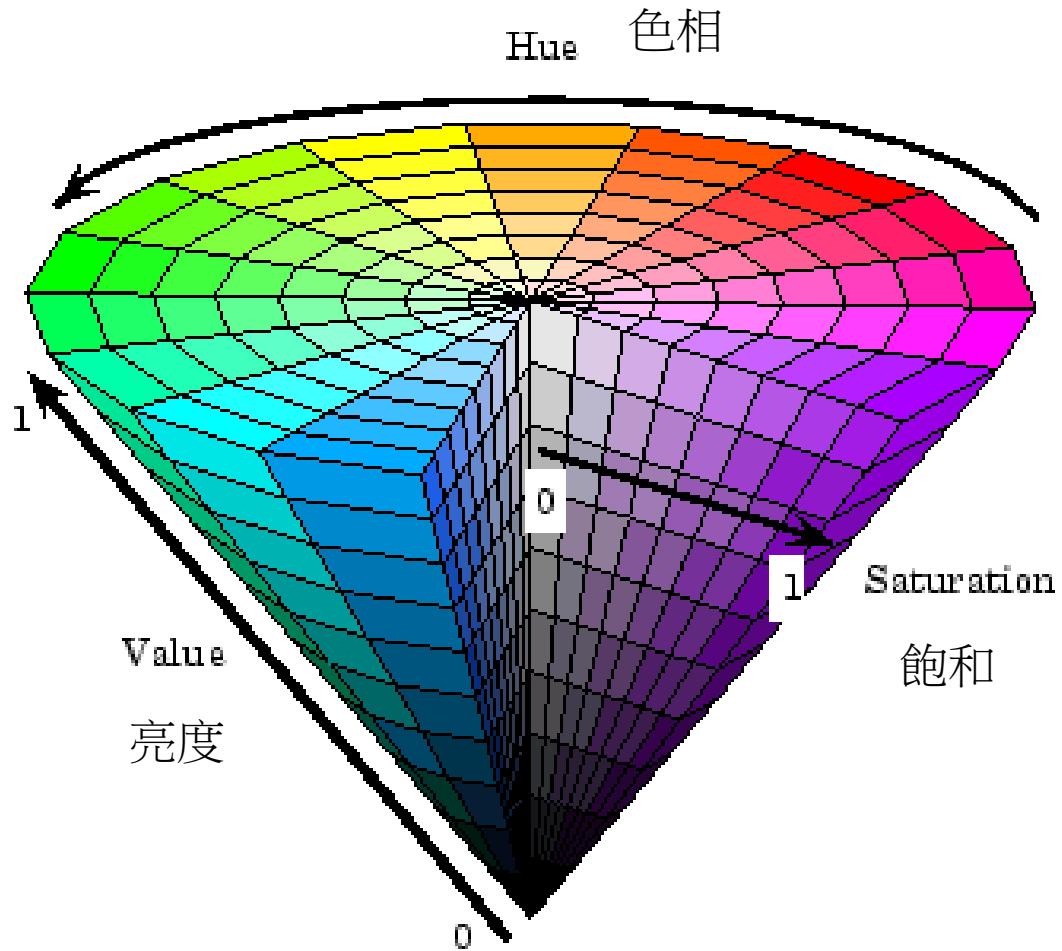
- Strongly correlated channels
- Non-perceptual

Image from: [http://en.wikipedia.org/wiki/File:RGB\\_color\\_solid\\_cube.png](http://en.wikipedia.org/wiki/File:RGB_color_solid_cube.png)

# Color spaces: HSV



## Intuitive color space



**H**  
( $S=1, V=1$ )



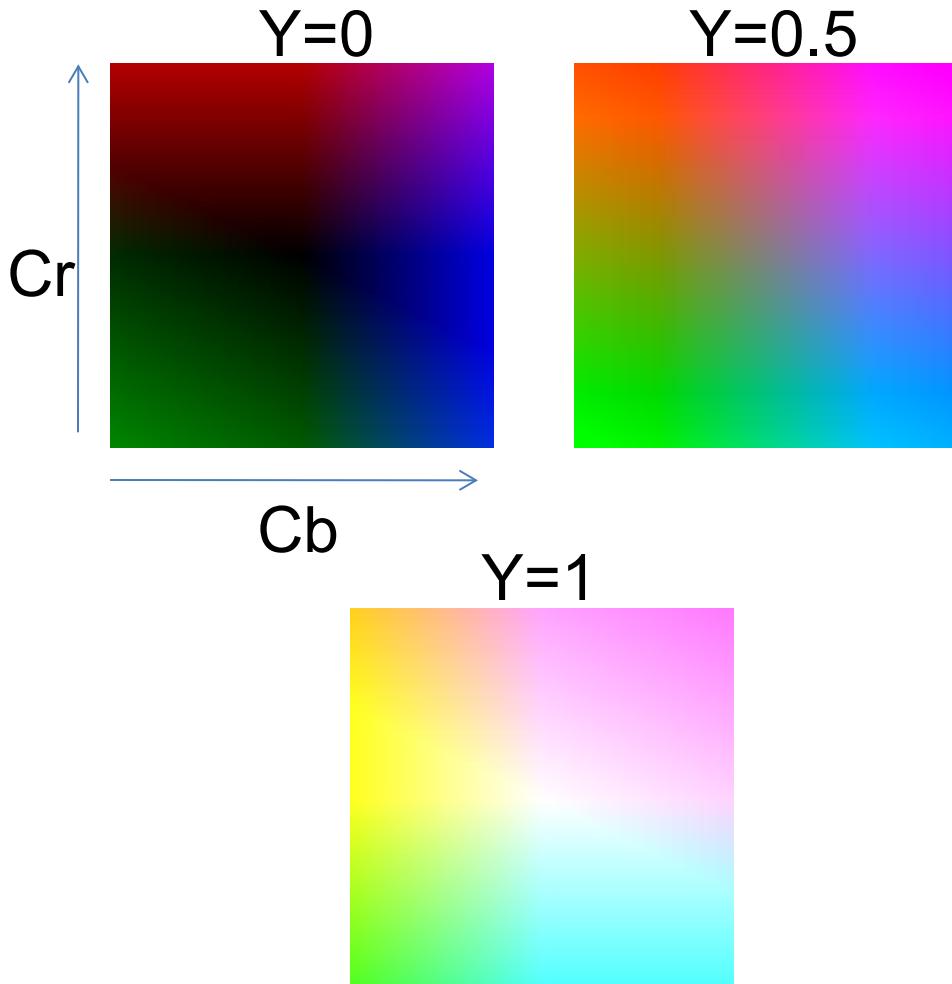
**S**  
( $H=1, V=1$ )



**V**  
( $H=1, S=0$ )

# Color spaces: YCbCr

Fast to compute, good for compression, used by TV



**Y**  
(Cb=0.5,Cr=0.5)



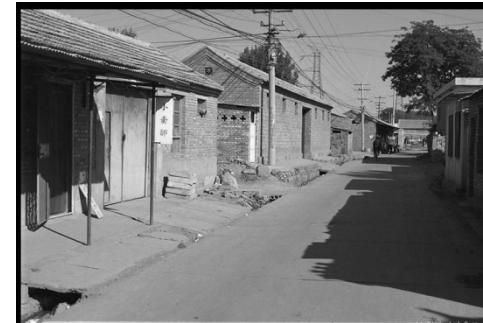
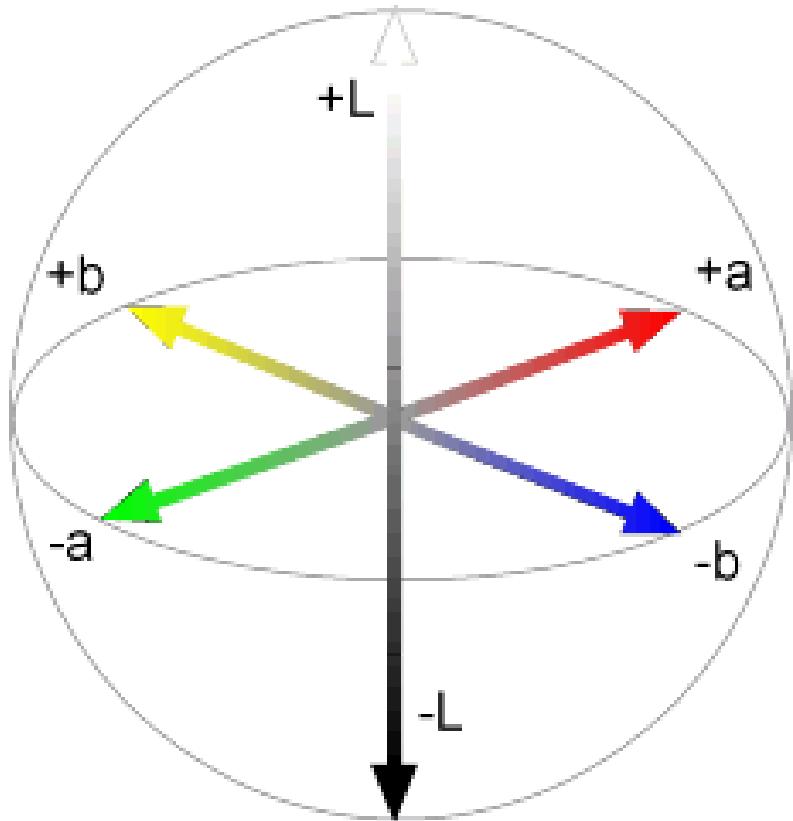
**Cb**  
(Y=0.5,Cr=0.5)



**Cr**  
(Y=0.5,Cb=0.5)

# Color spaces: L\*a\*b\*

“Perceptually uniform”\* color space



**L**  
( $a=0, b=0$ )



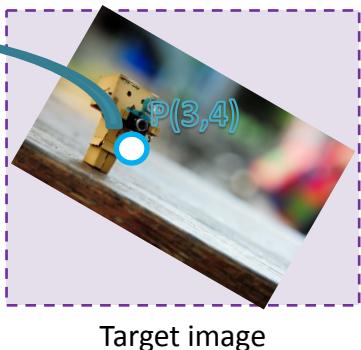
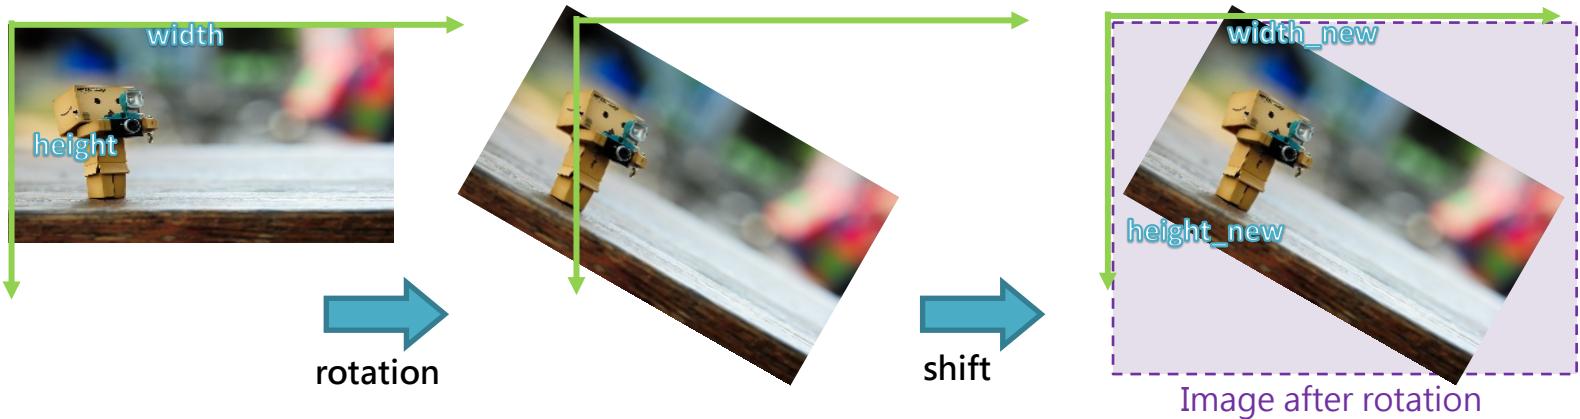
**a**  
( $L=65, b=0$ )



**b**  
( $L=65, a=0$ )

# Example: Rotation (Lab 5)

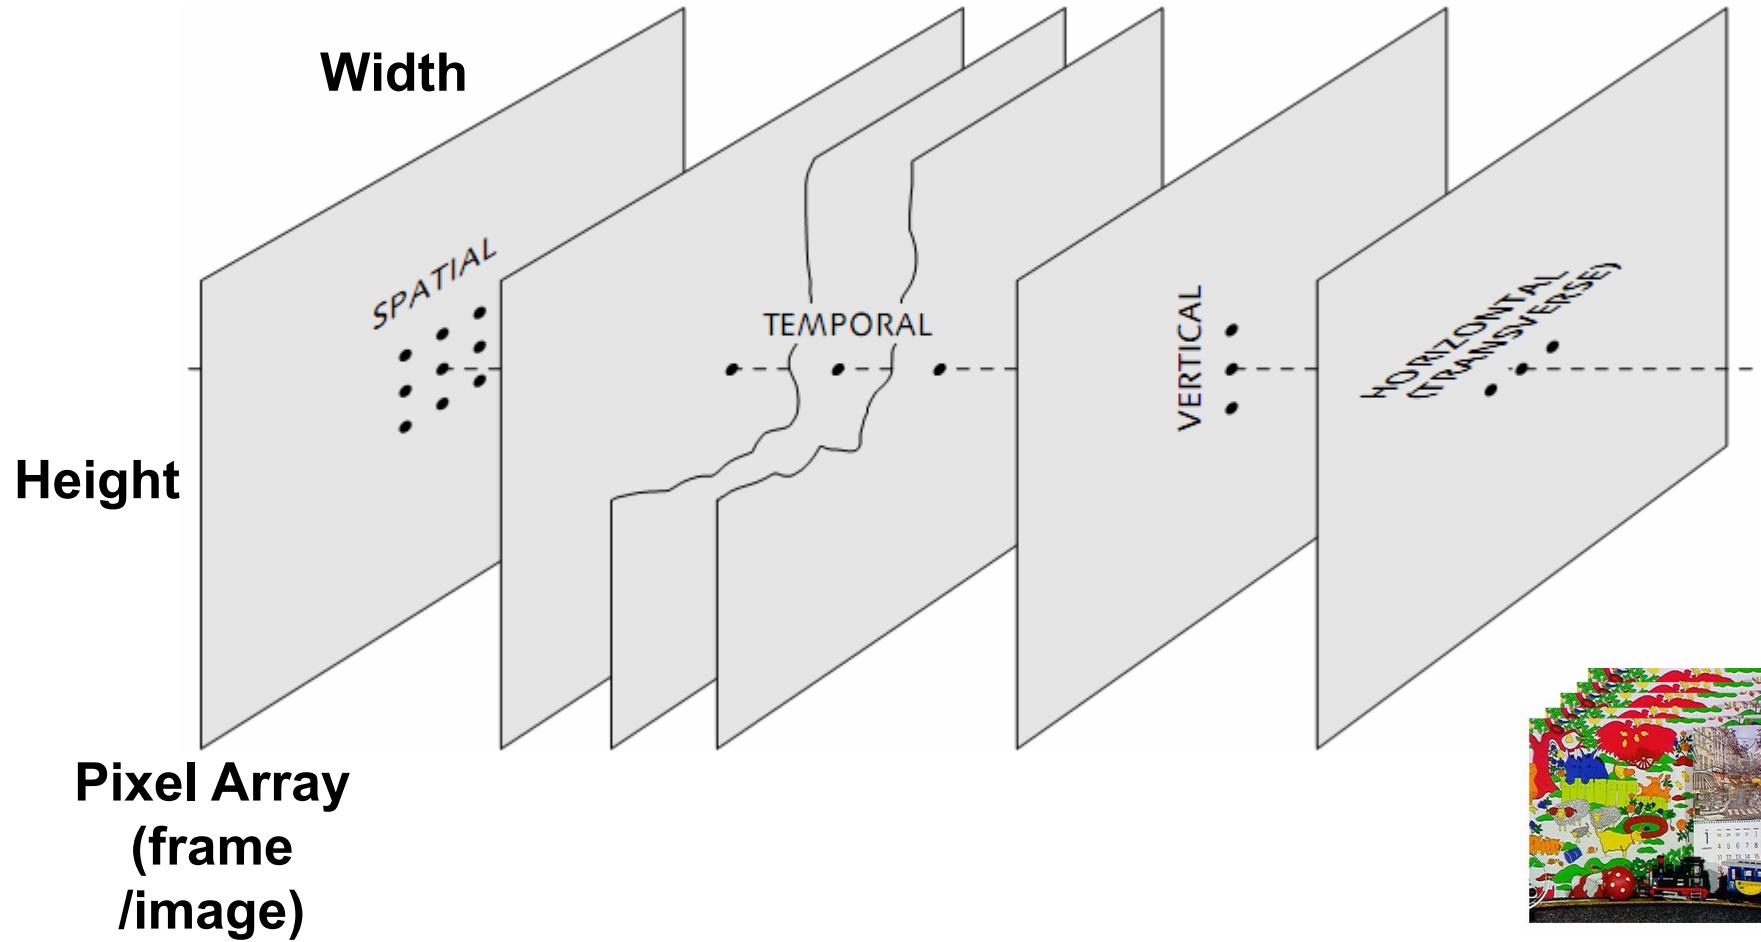
(1,1)



e.g.  
Backward warping +  
bilinear interpolation

# Video = Sequence of Images

3-D index: (x, y, t)



Ref: *Digital Video and HDTV*



# Refresh Rate v.s. Frame Rate

<i>Viewing environment</i>	<i>Ambient illumination</i>	<i>Refresh (flash) rate, Hz</i>	<i>Frame rate, Hz</i>
Cinema	Dark	48	24
Television	Dim	50	25
	Dim	≈60	≈30
Office	Bright	various, e.g., 66, 72, 76, 85	same as refresh rate

At least high enough to avoid flicker. ↪

Ref: *Digital Video and HDTV*

Higher for fast-moving objects.

**Ex. for 120Hz LCD TV:**

**Encoded Video: 30Hz (frame per second, fps)**

**HDMI 1080p : 60Hz (progressive)**

**Refresh : 120Hz**

# Resolution Explosion



**85" 8Kx4K (sharp)**



**84" 4Kx2K (LG)**

monitor

(Eizo, ViewSonic,  
IBM, Sharp, ...)

**20"~36" 4Kx2K**



**3.5" 320p  
(Apple)**



**3.5" 640p  
(Apple)**



**4" 640p  
(Apple)    5" 1080p  
(HTC)**



**2009**

**2010**

**2011**

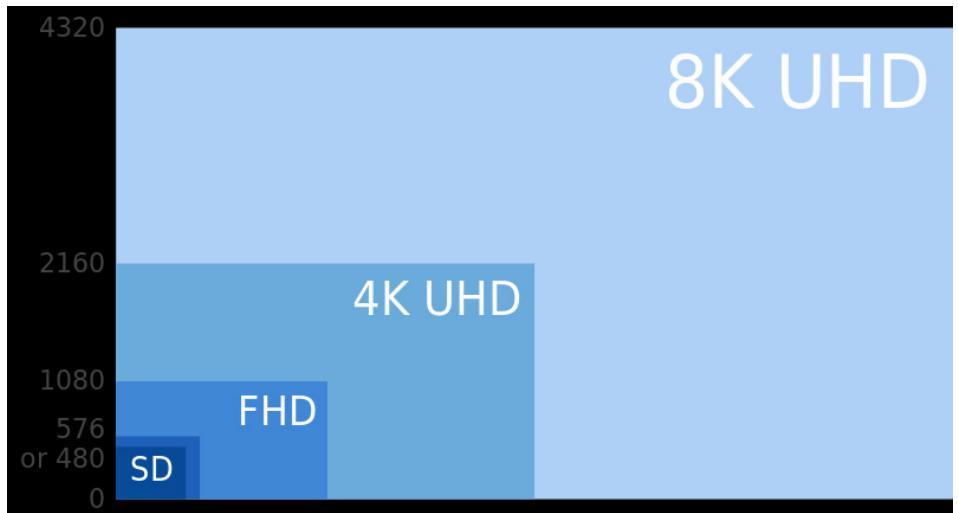
**2012**



**1080p  
iPhone6  
2014**



# Frame Rate Going Up



**Ultra High Definition TV**  
4Kx2K/8Kx4K  
Up to 120 fps by 2020



**Cinema**  
HFR (High Frame Rate)  
48fps



**LCD/LED TV**  
**240Hz** refresh rate



# Audio/Image/Video in Practice

- Audio
  - 1-D signal
  - Intensity
  - Long-tap filter
  - Throughput
    - ~ 44K samples/s
- Image
  - 2-D signal
  - RGB triple
  - Short-tap filter
    - Locally stationary
  - Large size
    - e.g. 24M pixels
- Video
  - 2-D + time
  - YCbCr triple
  - Motion model
  - Resource demanding
    - 8K4K 120fps  
=> 3.8Gpixels/s
    - > 60% of internet traffic



# Lecture Outline

- Image and Video Basics
- Recap on Image filtering
- Interest Point
- Seam Carving



# Image filtering

- Image filtering: compute function of local neighborhood at each position
- Really important!
  - Enhance images
    - Denoise, resize, increase contrast, etc.
  - Extract information from images
    - Texture, edges, distinctive points, etc. **[Lab]**
  - Detect patterns
    - Template matching

# Example: box filter

$g[\cdot, \cdot]$

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

# Image filtering



$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

 $f[., .]$  $h[., .]$ 

0

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

0										

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

# Image filtering



$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

 $f[., .]$  $h[., .]$ 

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

0			10							

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

# Image filtering



$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

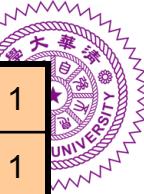
 $f[., .]$  $h[., .]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	0	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

			0	10	20				

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

# Image filtering



$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

 $f[., .]$  $h[., .]$ 

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

# Image filtering

$$g[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

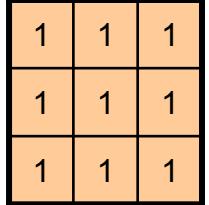
$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$


$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$$f[., .]$$

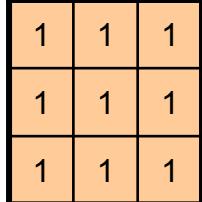
$$h[., .]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

			0	10	20	30	30		

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

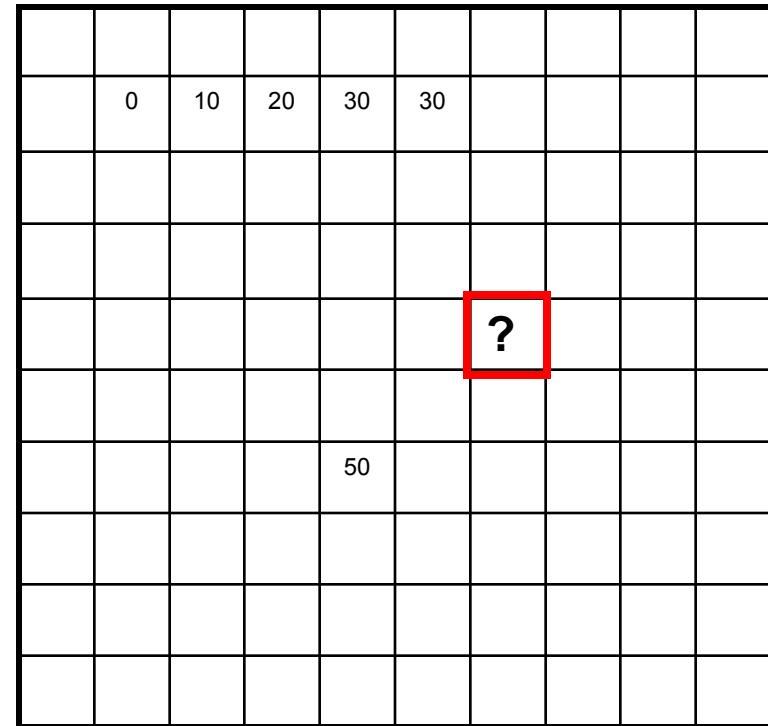
# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$$f[., .]$$

$$h[., .]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

# Image filtering

$$g[\cdot, \cdot] \quad \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

# Box Filter

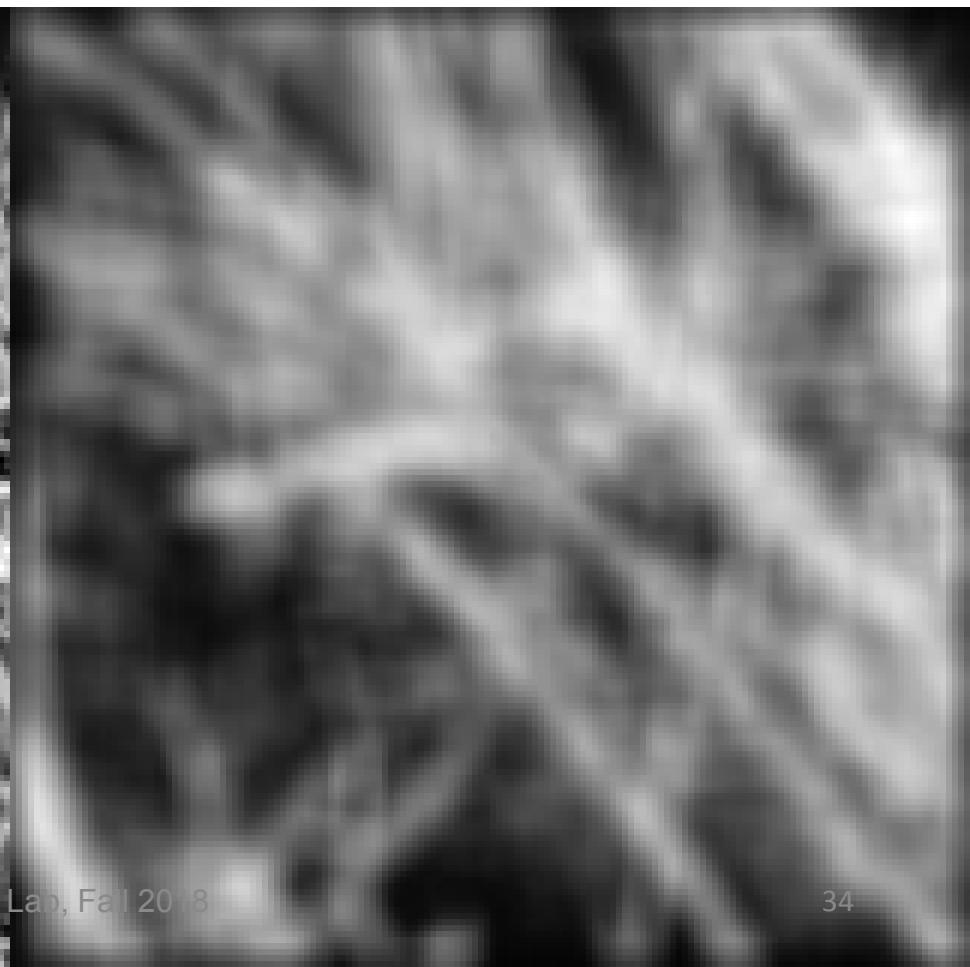
What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

# Smoothing with box filter



# Practice with linear filters



0	0	0
0	1	0
0	0	0

?

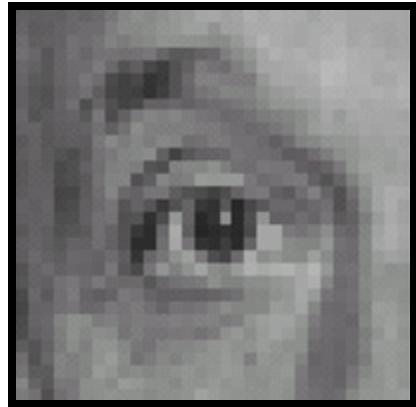
Original

# Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)



# Practice with linear filters



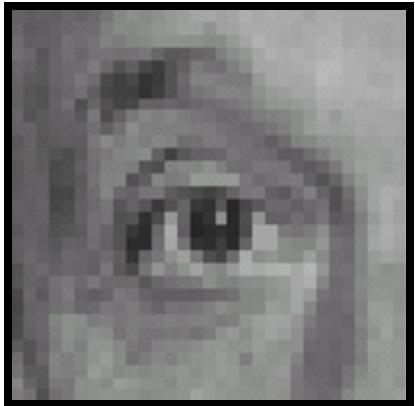
Original

0	0	0
0	0	1
0	0	0

?

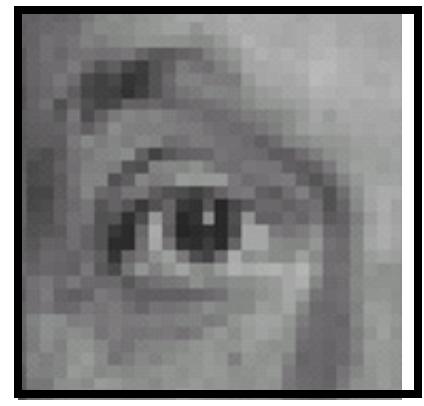


# Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

# Practice with linear filters



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

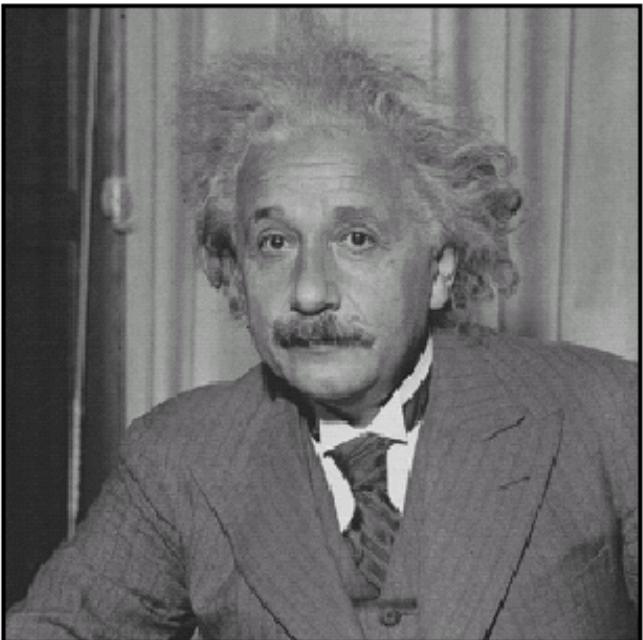


Original

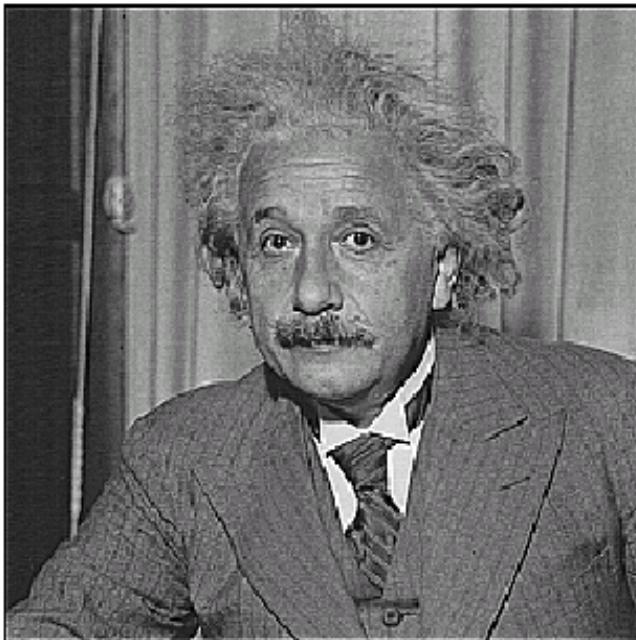
## Sharpening filter

- Accentuates (突顯) differences with local average

# Sharpening

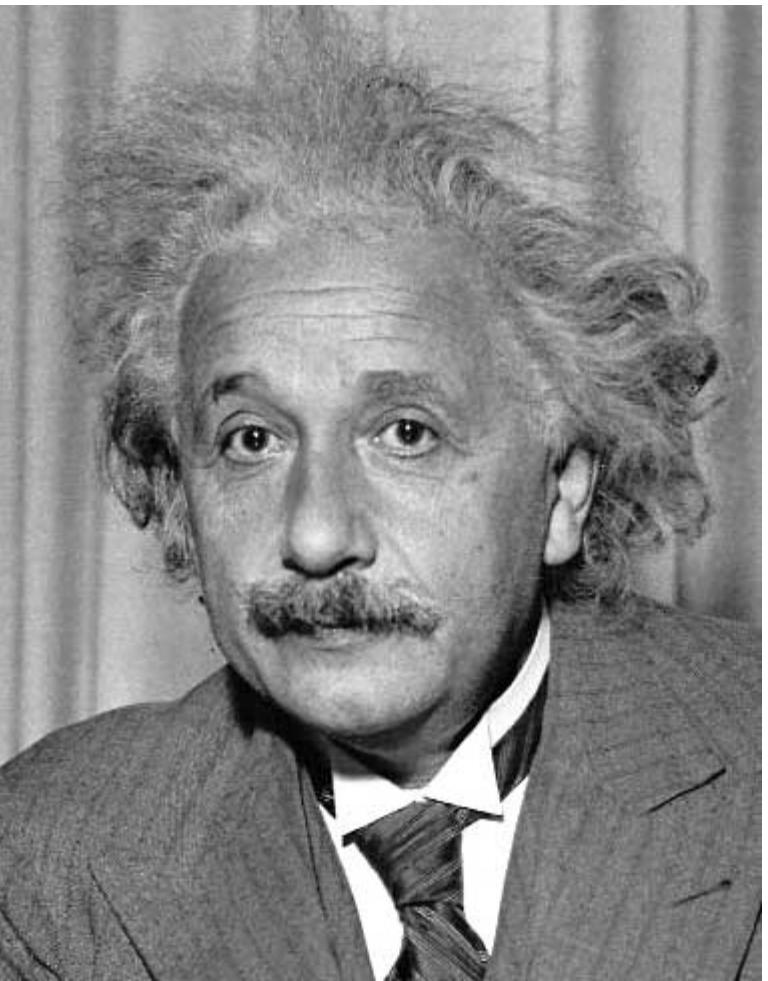


**before**



**after**

# Other filters



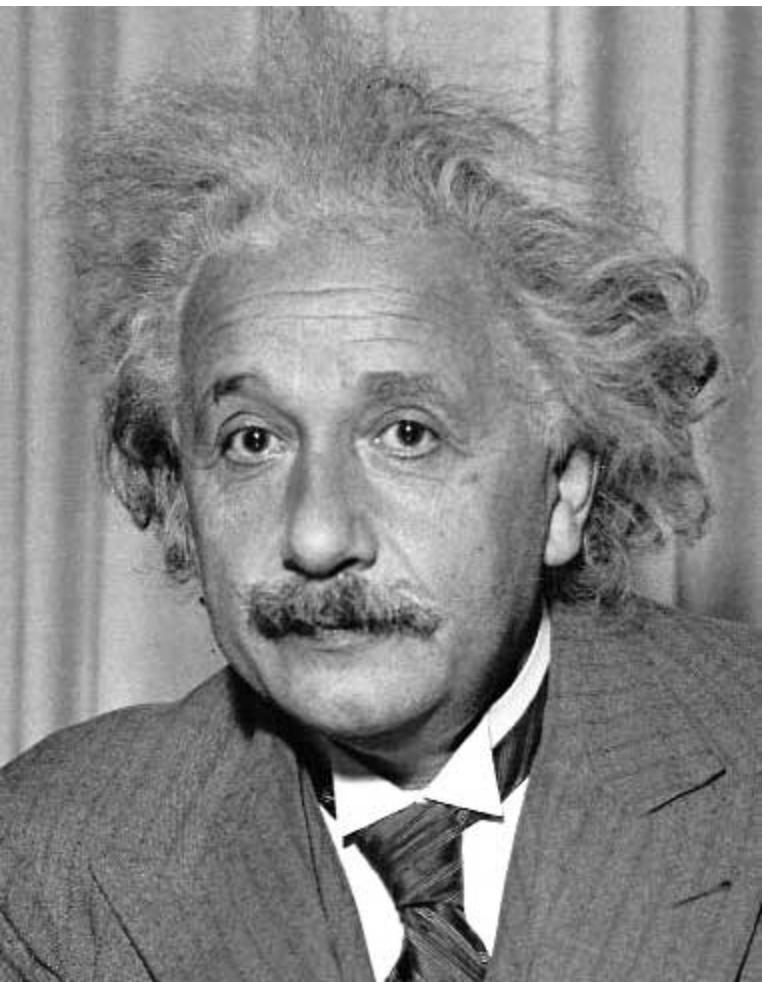
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge  
(absolute value)

# Other filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge  
(absolute value)



# Filtering vs. Convolution

- filtering

- $h = \text{filter2}(g, f)$ ; or

- $h = \text{imfilter}(f, g)$ ;

$$h[m, n] = \sum_{k, l} g[k, l] f[m+k, n+l]$$

g=filter      f=image

- 2d convolution

- $h = \text{conv2}(g, f)$ ;

$$h[m, n] = \sum_{k, l} g[k, l] f[m-k, n-l]$$

$$h[m, n] = \sum_{k, l} g[-k, -l] f[m+k, n+l]$$



# Key properties of linear filters

## Linearity:

$$\text{Conv2}(g_1 + g_2, f) = \text{conv2}(g_1, f) + \text{conv2}(g_2, f)$$

**Shift invariance:** same behavior regardless of pixel location

$$\text{Conv2}(\text{shift}(g), f) = \text{shift}(\text{conv2}(g, f))$$

Any linear, shift-invariant operator can be represented as a **convolution**



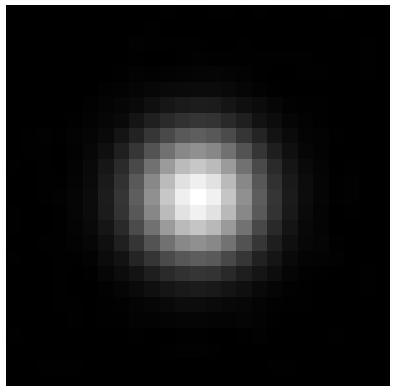
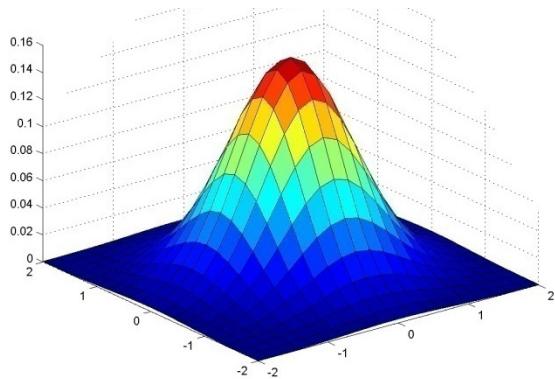
# More properties

- Commutative:  $a * b = b * a$ 
  - Conceptually no difference between filter and signal
  - But particular filtering implementations might break this equality
- Associative:  $a * (b * c) = (a * b) * c$ 
  - Often apply several filters one after another:  $((a * b_1) * b_2) * b_3$
  - This is equivalent to applying one filter:  $a * (b_1 * b_2 * b_3)$
- Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out:  $ka * b = a * kb = k(a * b)$
- Identity: unit impulse  $e = [0, 0, 1, 0, 0]$ ,  
 $a * e = a$

# Important filter: Gaussian



- Weight contributions of neighboring pixels by nearness



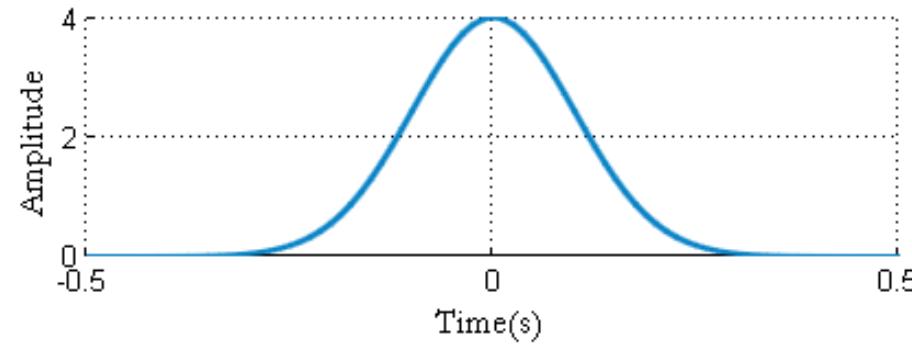
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

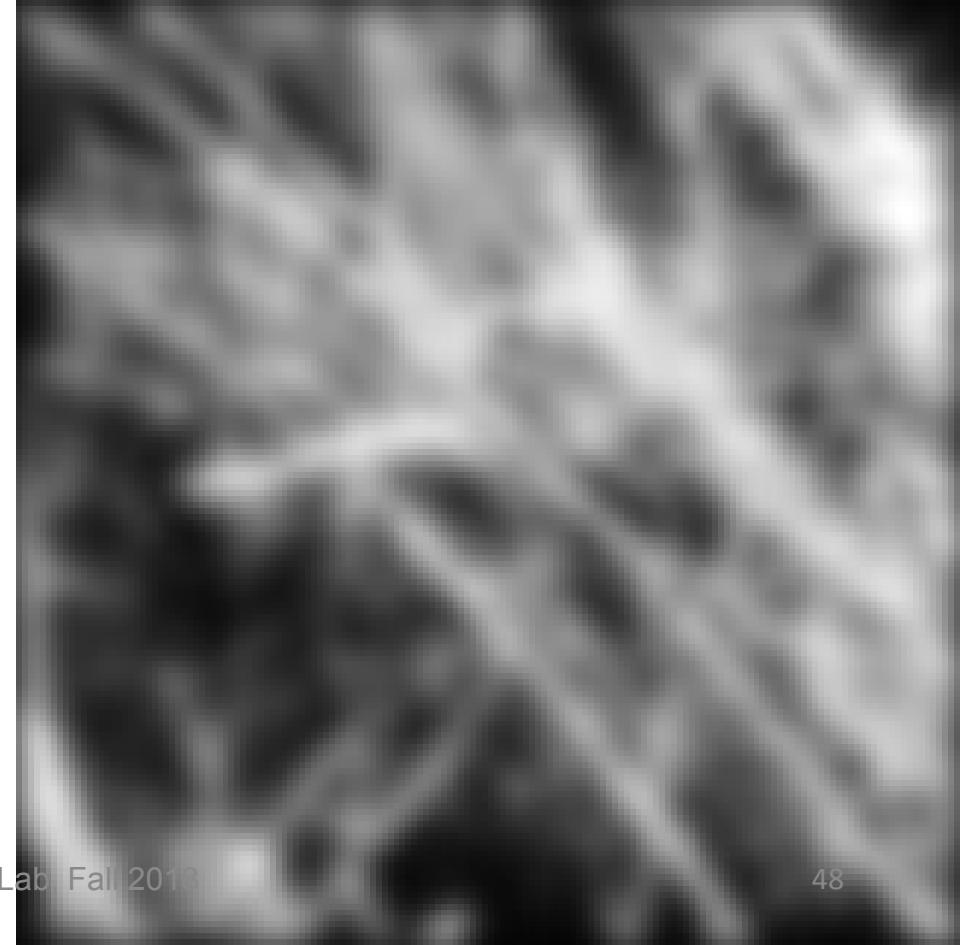
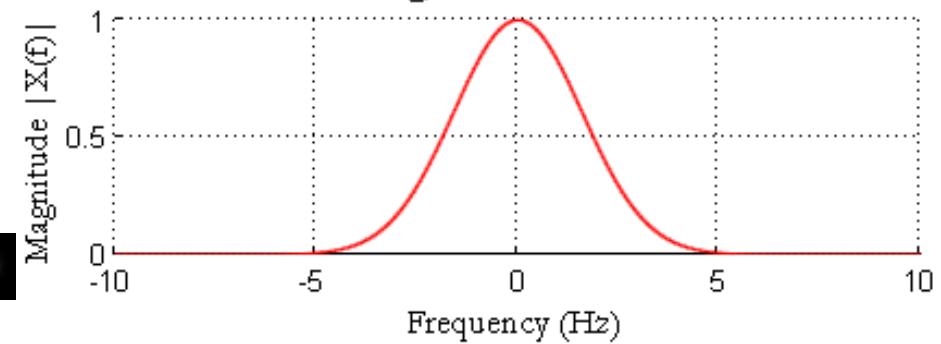
$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Smoothing with Gaussian filter

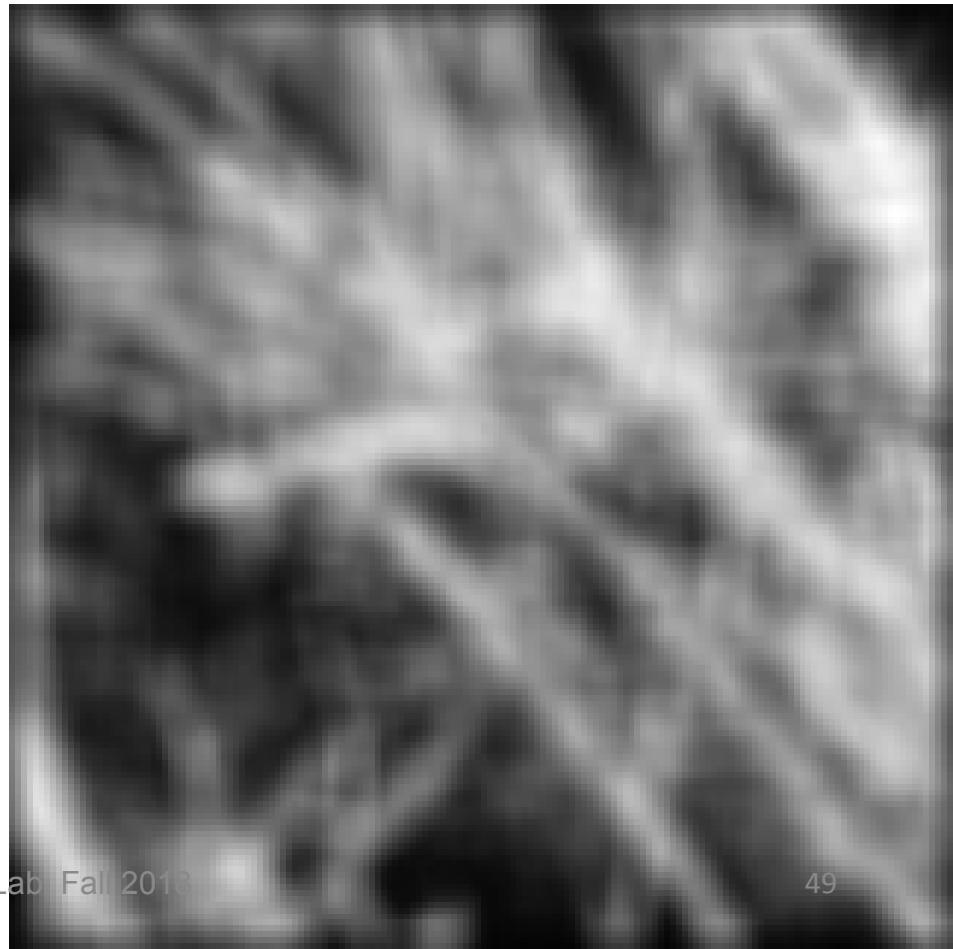
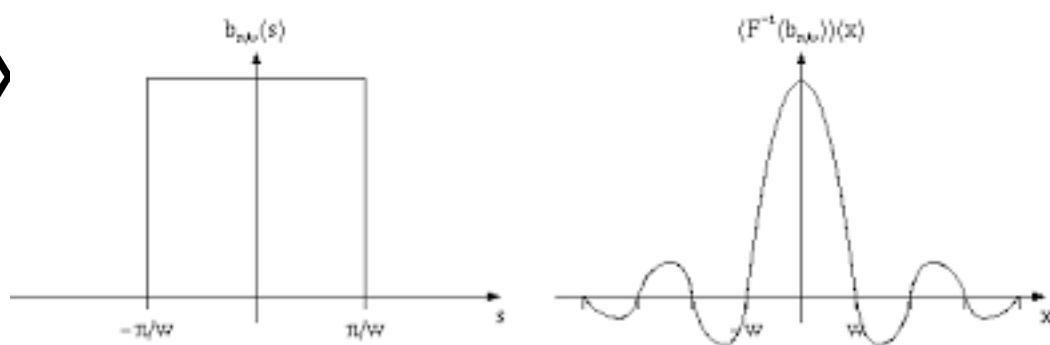
Gaussian Pulse  $\sigma=0.1s$



Magnitude of FFT



# Smoothing with boxcar





# Gaussian filters

- Remove “high-frequency” components from the image (**low-pass filter**)
  - Images become more smooth
- Convolution with self is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable kernel*
  - Factors into product of two 1D Gaussians



# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian



# Separability example

2D convolution  
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform convolution  
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 \\ 18 \\ 18 \end{matrix}$$

Followed by convolution  
along the remaining column:



# Separability

- Why is separability useful in practice?
  - Full filter: Filtering an  $M$ -by- $N$  image with a  $P$ -by- $Q$  filter kernel requires roughly  $MNPQ$  multiplies and adds.
  - Separable filter: The first step requires about  $MNP$  multiplies and adds. The second requires about  $MNQ$  multiplies and adds, for a total of  $MN(P + Q)$ .

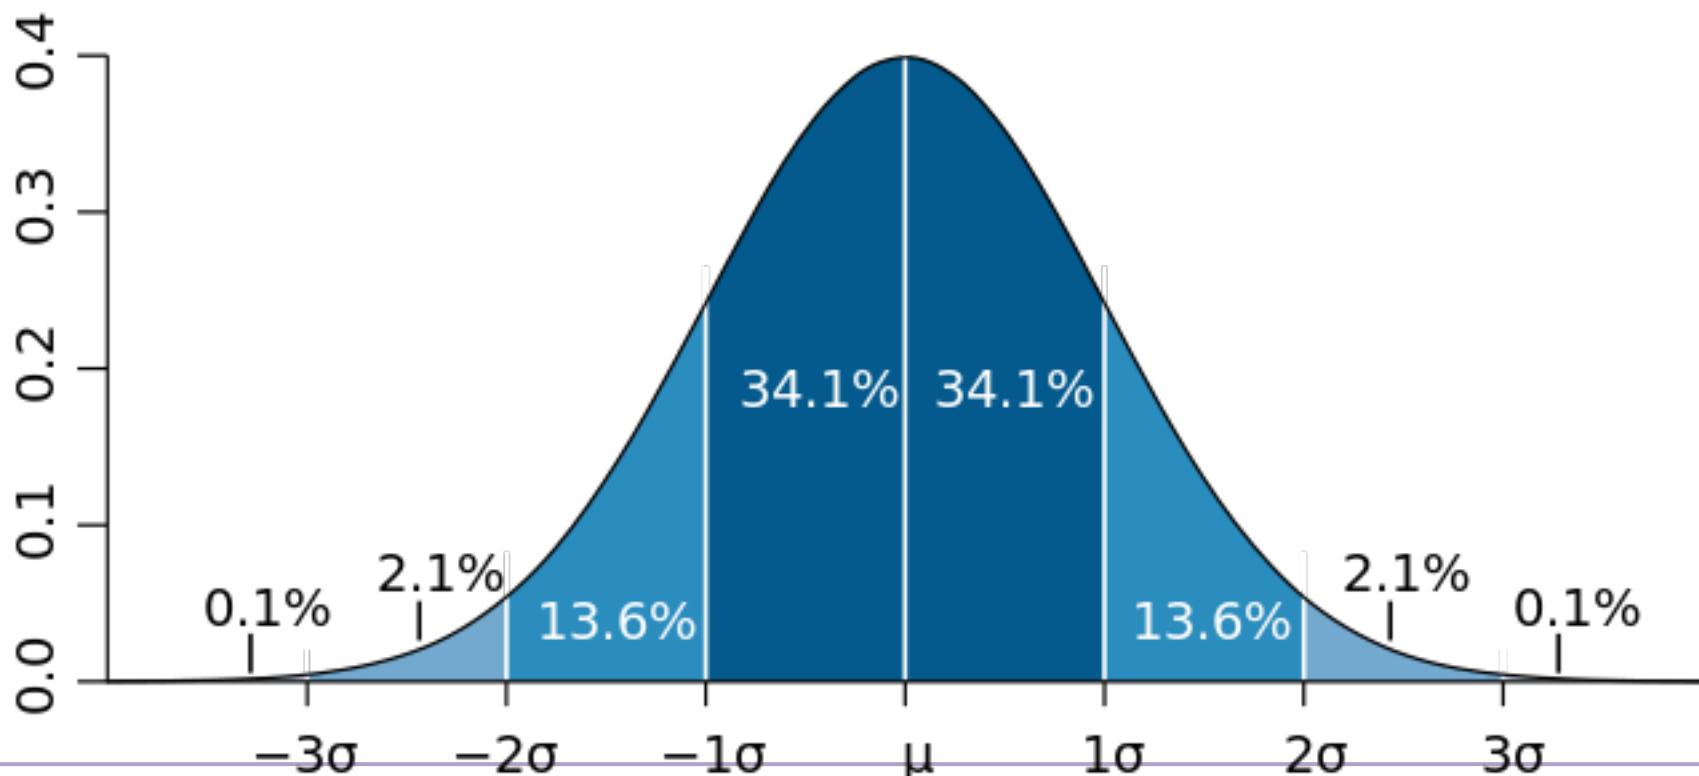
<http://blogs.mathworks.com/steve/2006/10/04/separable-convolution/>

# Practical matters



## How big should the filter be?

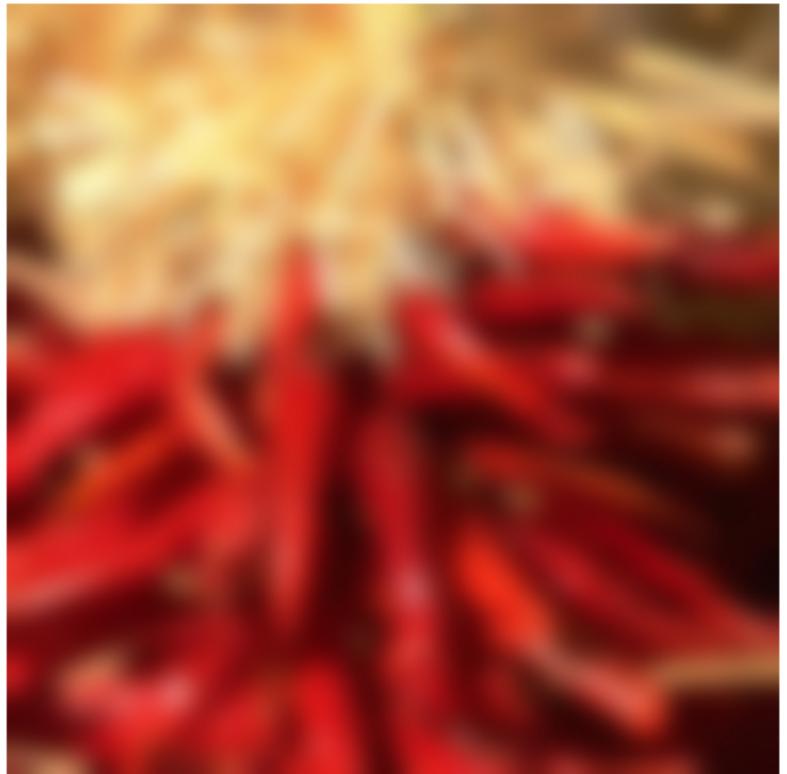
- Values at edges should be near zero
- Rule of thumb for Gaussian: set filter half-width to about  $3\sigma$





# Practical matters

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge



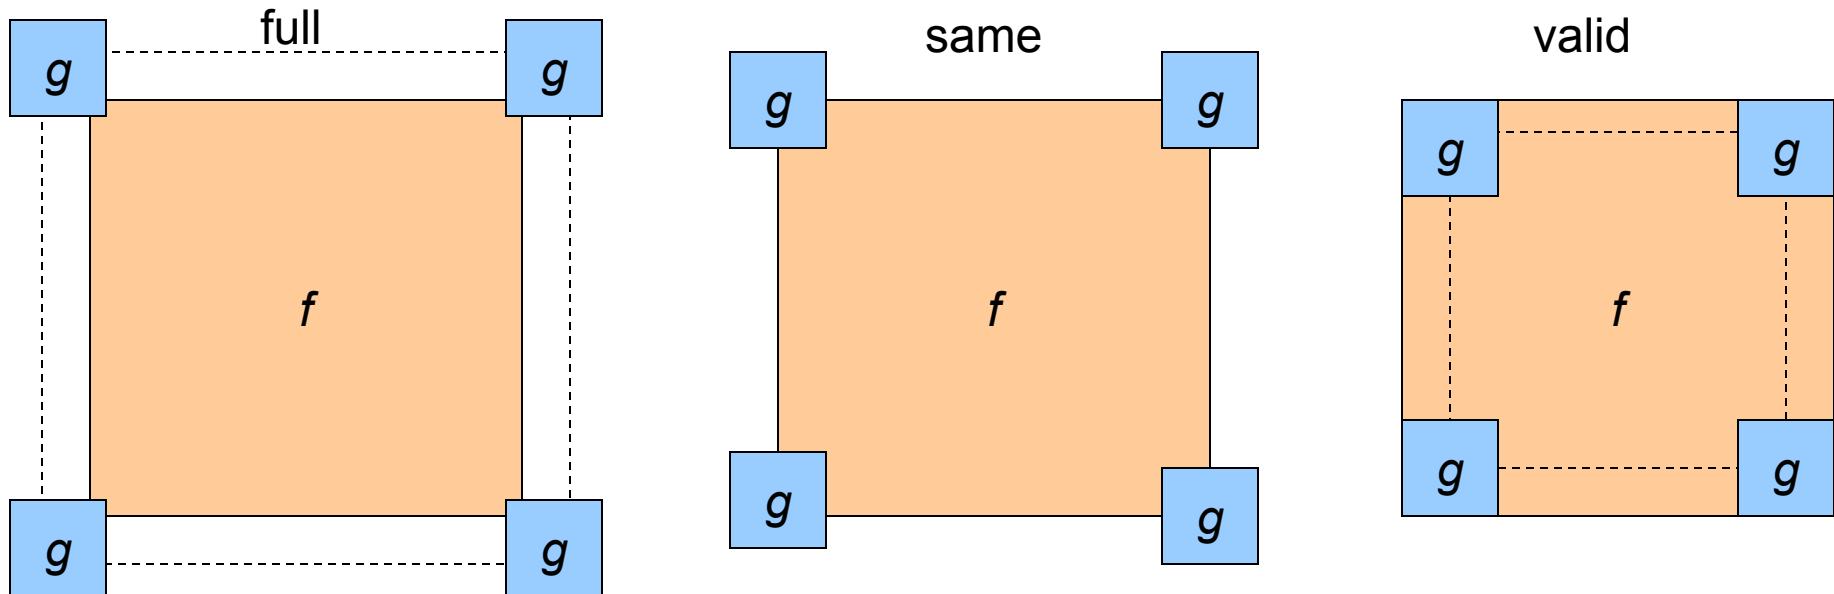


# Practical matters

- methods (MATLAB):
  - clip filter (black):      `imfilter(f, g, 0)`
  - wrap around:            `imfilter(f, g, 'circular')`
  - copy edge:                `imfilter(f, g, 'replicate')`
  - reflect across edge: `imfilter(f, g, 'symmetric')`

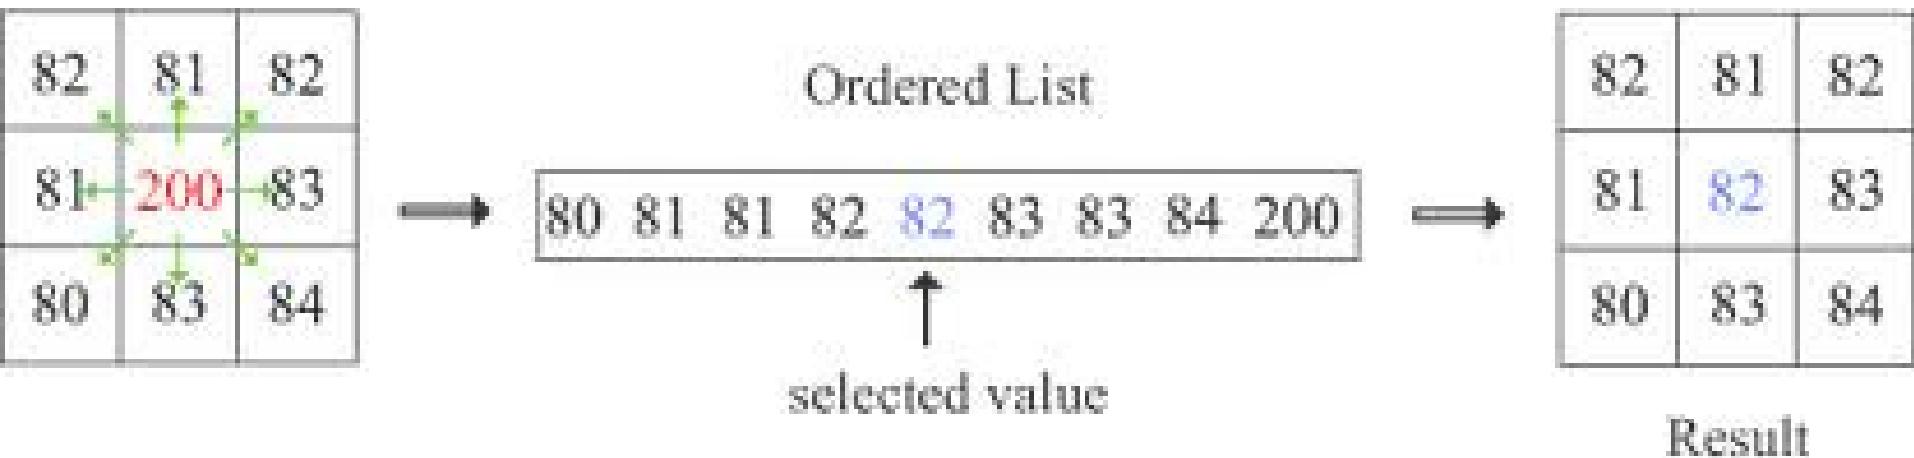
# Practical matters

- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
  - `shape = 'full'` : output size is sum of sizes of  $f$  and  $g$
  - `shape = 'same'` : output size is same as  $f$
  - `shape = 'valid'` : output size is difference of sizes of  $f$  and  $g$



# Median filters

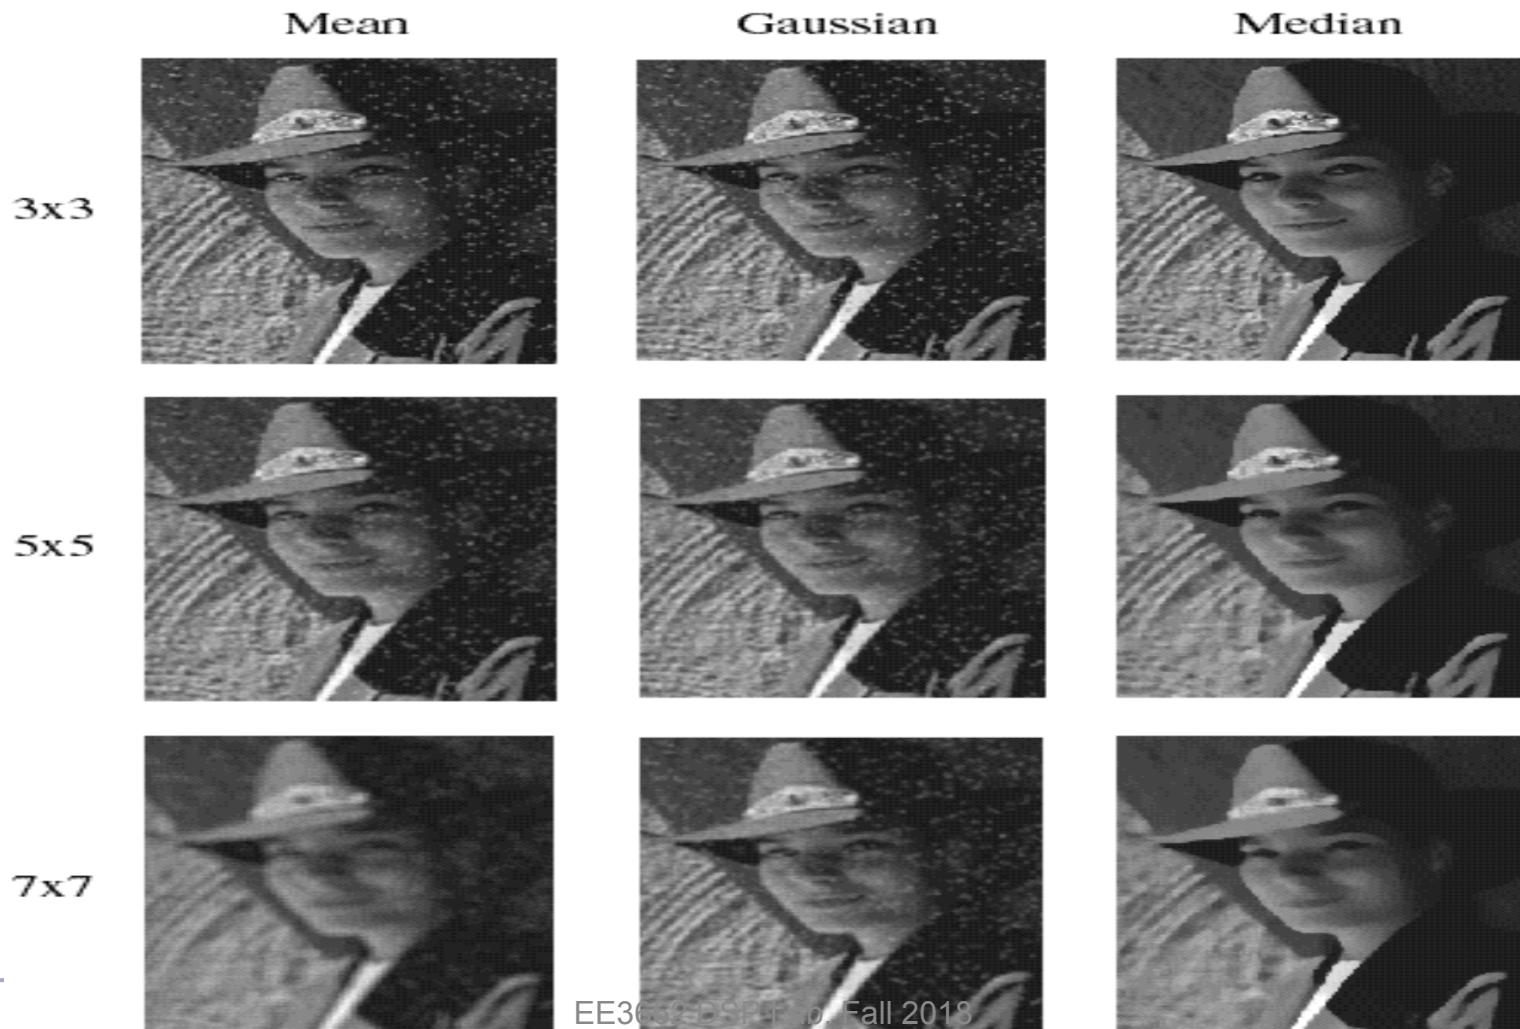
- A **Median Filter** operates over a window by selecting the median intensity in the window.



$B = \text{ordfilt2}(A, \text{order}, \text{domain})$   
In matlab

# Median filters

- What advantage does a median filter have over a mean filter? (salt and pepper noise)



# Median filters

- Is a median filter a kind of convolution?

$$h[m, n] = \sum_{k,l} g[k, l] f[m+k, n+l]$$

$$h[m, n] = \text{median}_{kl}(f[m+k, n+l])$$

1	1	1
2	1	1
1	1	1

1	1	1
5	1	1
1	1	1

1	5	1
1	1	1
1	1	1

Median:

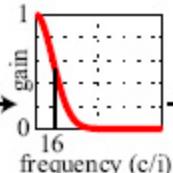
1

1

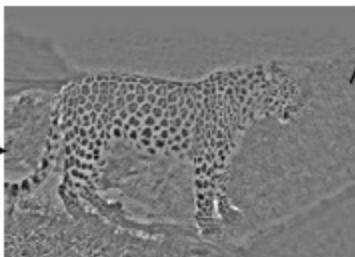
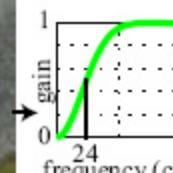
1

# Lab 6: Hybrid Images

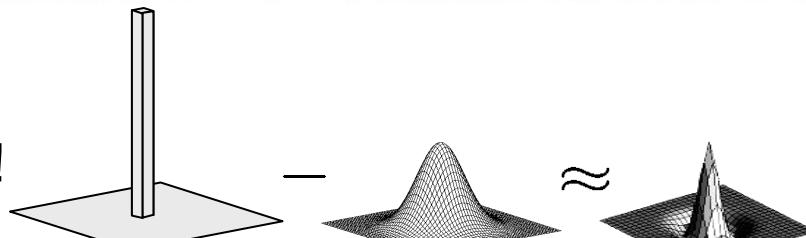
Gaussian Filter!



A. Oliva, A. Torralba, P.G. Schyns,  
“Hybrid Images,” SIGGRAPH 2006



Laplacian Filter!



unit impulse

Gaussian

Laplacian of Gaussian

**Gaussian Filter  
+Sampling**





# Lecture Outline

- Image and Video Basics
- Recap on Image Filtering
- Interest Point
  - Corners
- Seam Carving



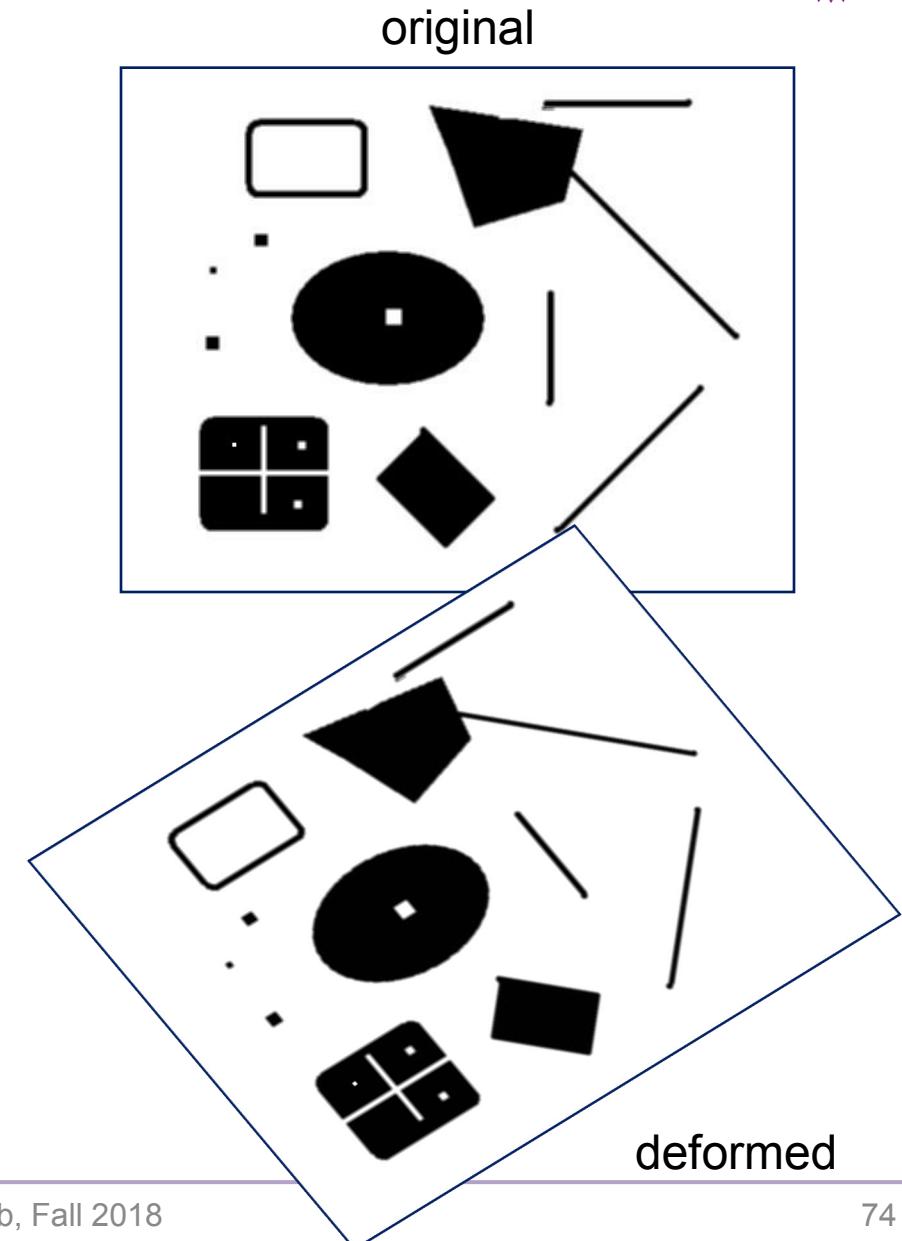
# Interest points

- Note: “interest points” = “keypoints”, also sometimes called “features”
- Many applications
  - tracking: which points are good to track?
  - recognition: find patches likely to tell us something about object category
  - 3D reconstruction: find correspondences across different views

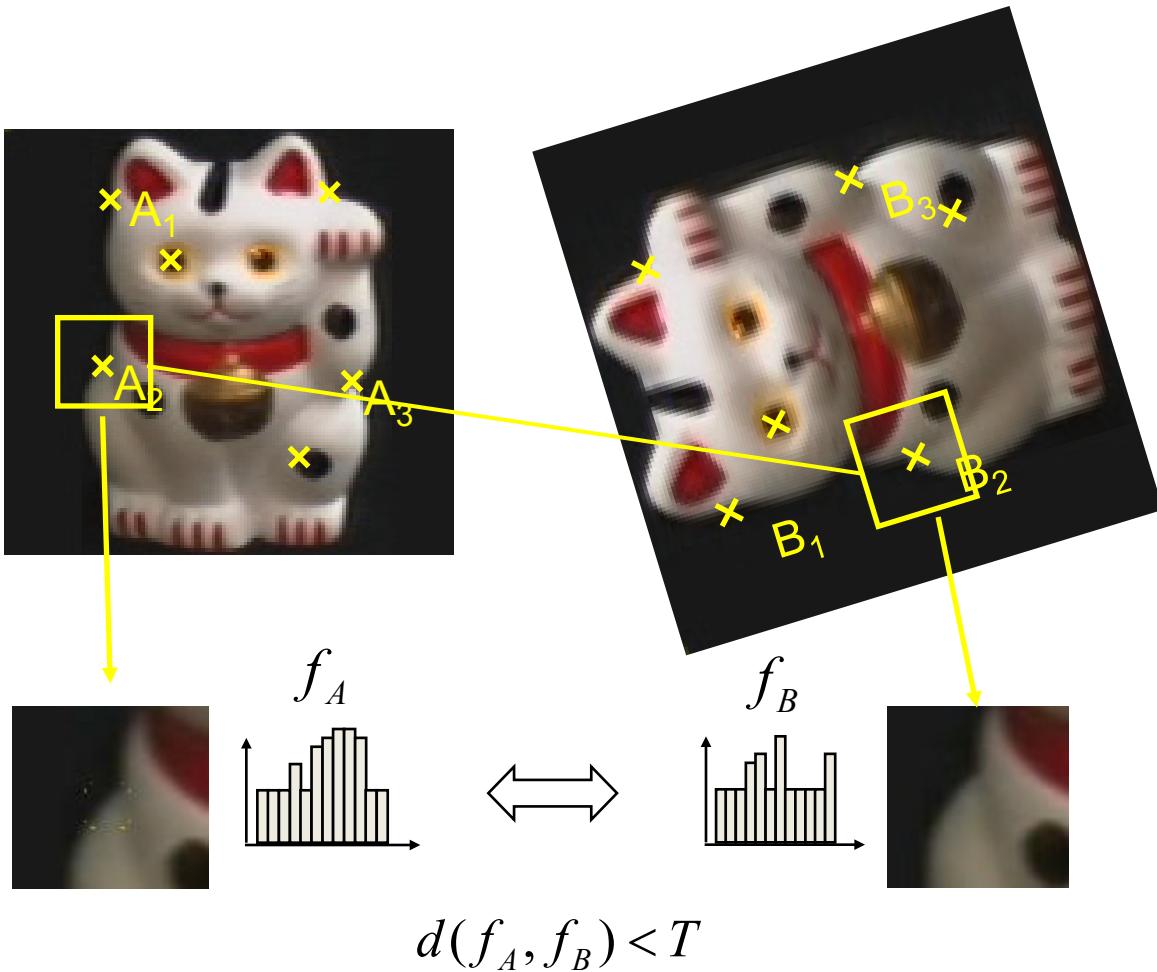


# Interest points

- Suppose you have to click on some point, go away and come back after I deform the image, and click on the same points again.
  - Which points would you choose?



# Overview of Keypoint Matching



- Find a set of distinctive keypoints**
- Define a region around each keypoint**
- Extract and normalize the region content**
- Compute a local descriptor from the normalized region**
- Match local descriptors**



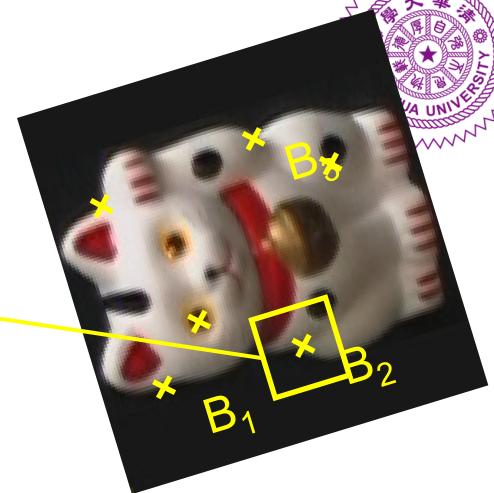
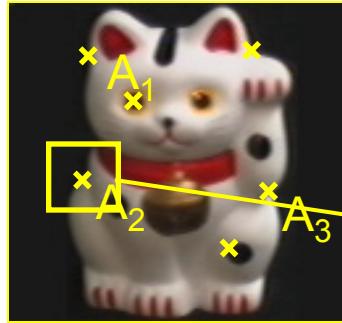
# Goals for Keypoints



Detect points that are *repeatable* and *distinctive*

*Repeatable*: the same feature can be found in several images despite geometric and photometric transformations

# Key trade-offs



## Detection of interest points



More Repeatable

Robust detection  
Precise localization

More Points

Robust to occlusion  
Works with less texture

## Description of patches



More Distinctive

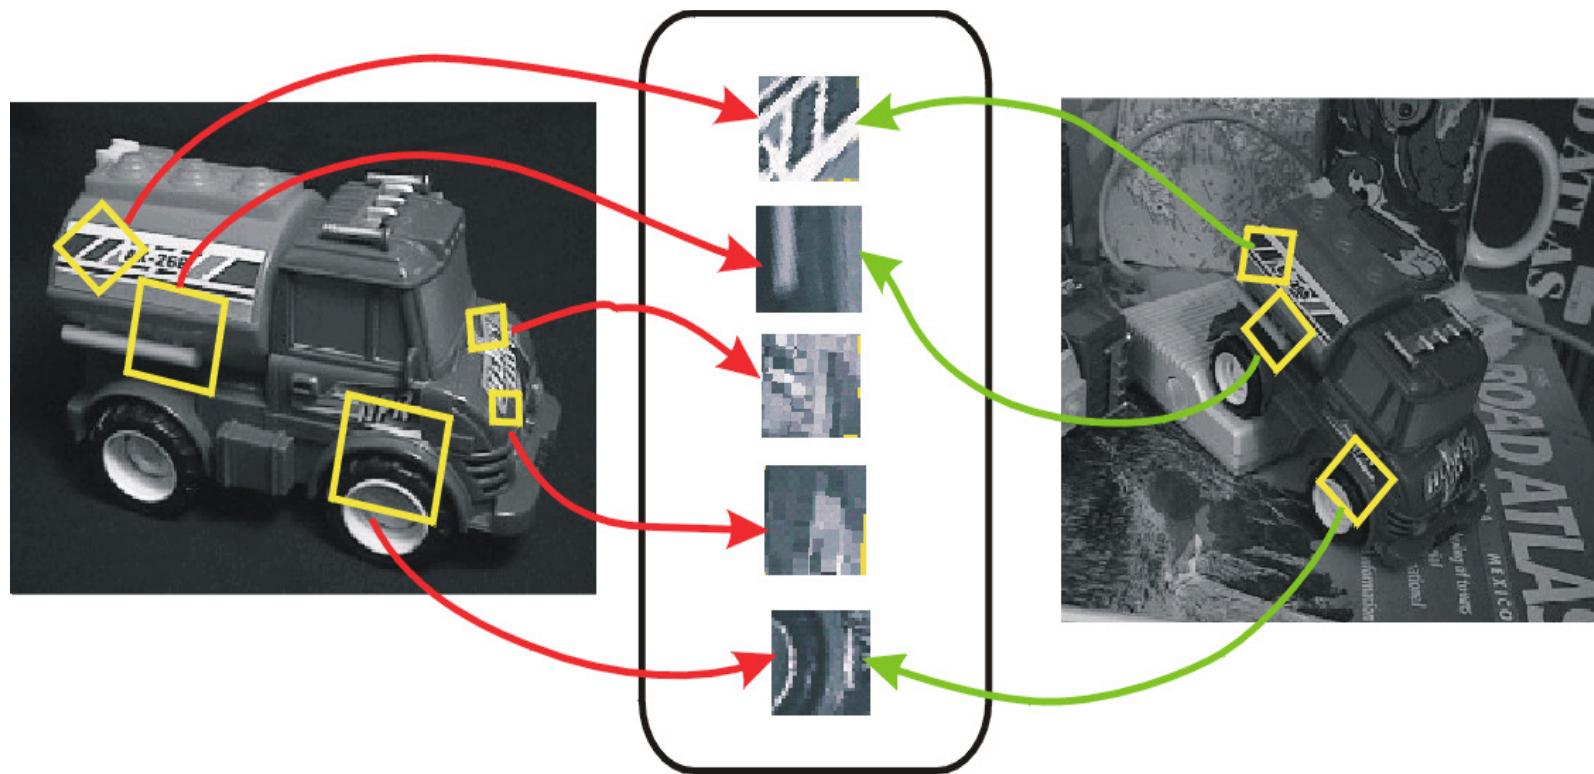
Minimize wrong matches

More Flexible

Robust to expected variations  
Maximize correct matches

# Invariant Local Features

- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



# Feature extraction: Corners

9300 Harris Corners Pkwy, Charlotte, NC





# Many Existing Detectors Available

Hessian & Harris

Laplacian, DoG

Harris-/Hessian-Laplace

Harris-/Hessian-Affine

EBR and IBR

MSER

Salient Regions

Others...

[Beaudet '78], [Harris '88]

[Lindeberg '98], [Lowe 1999]

[Mikolajczyk & Schmid '01]

[Mikolajczyk & Schmid '04]

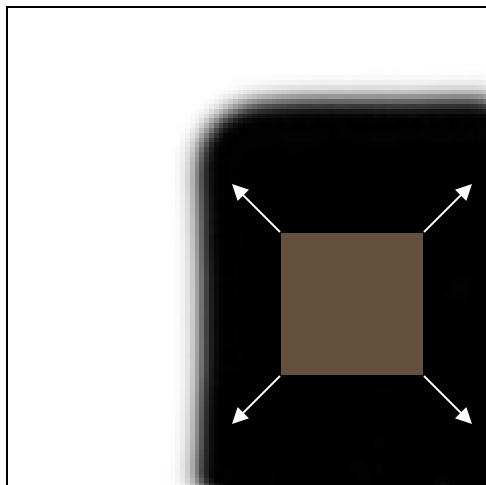
[Tuytelaars & Van Gool '04]

[Matas '02]

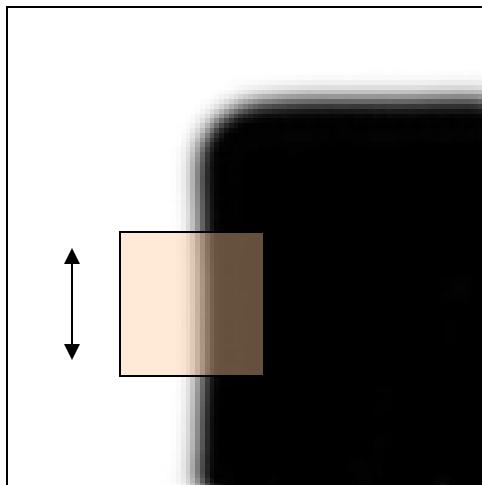
[Kadir & Brady '01]

# Corner Detection (Lab 7): Basic Idea

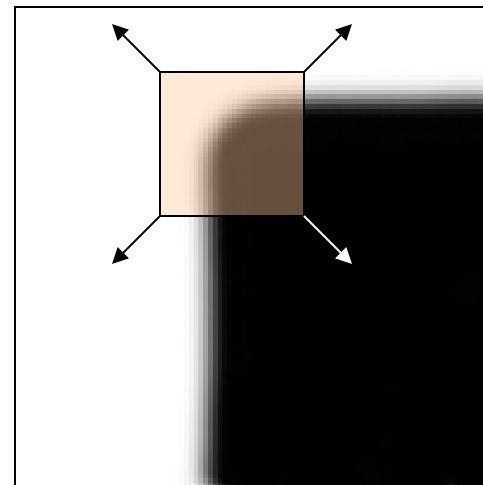
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



**“flat” region:**  
no change in  
all directions



**“edge”:**  
no change  
along the edge  
direction



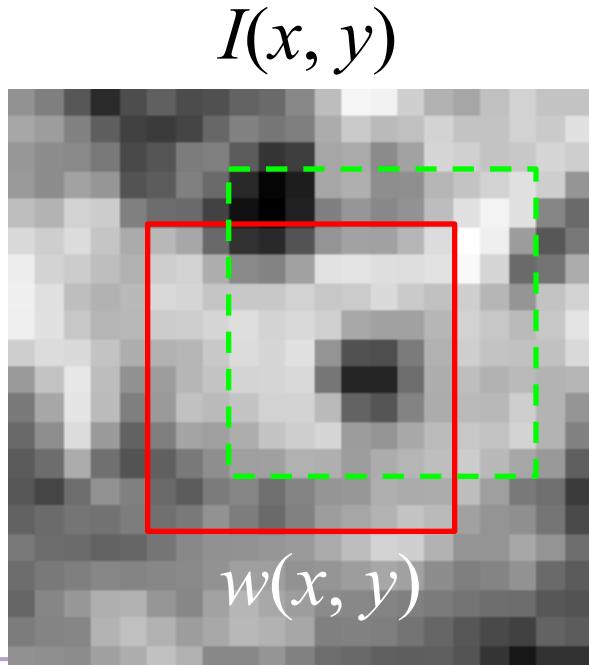
**“corner”:**  
significant  
change in all  
directions

# Corner Detection: Mathematics

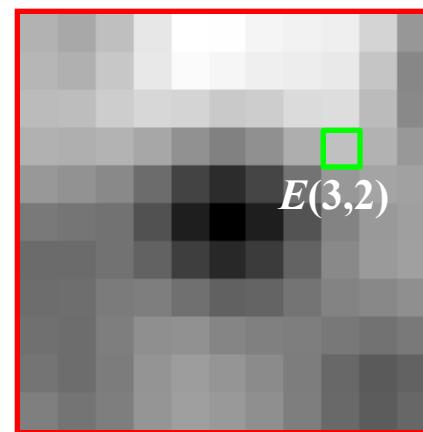


Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



$$E(u, v)$$

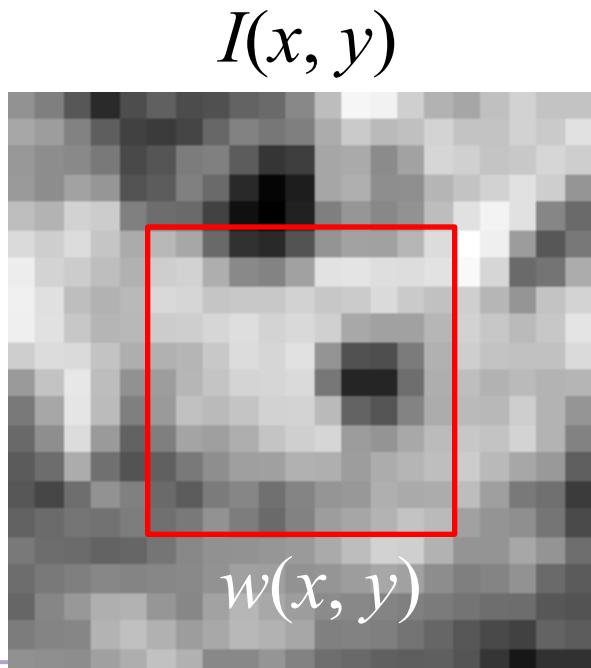


# Corner Detection: Mathematics

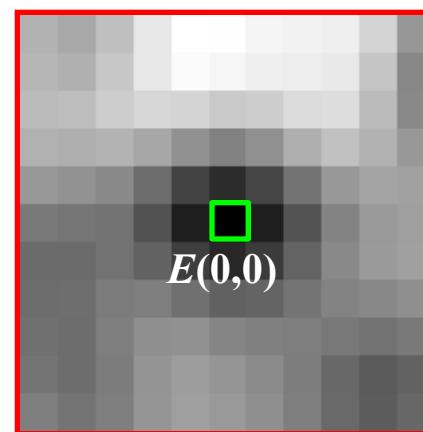


Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



$E(u, v)$



# Corner Detection: Mathematics

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

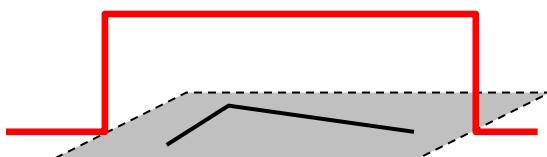
$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

Window  
function

Shifted  
intensity

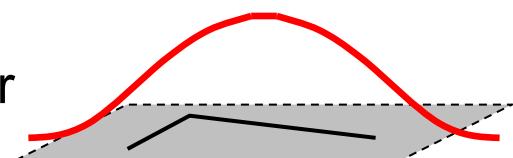
Intensity

Window function  $w(x,y) =$



1 in window, 0 outside

or



Gaussian



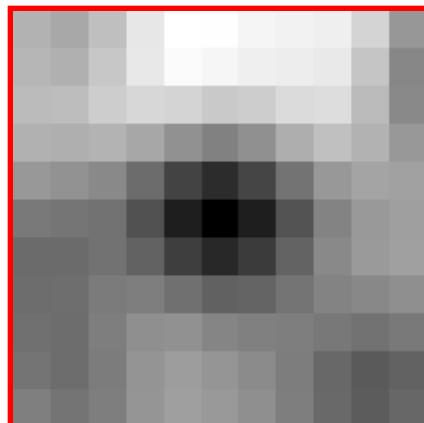
# Corner Detection: Mathematics

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$





# Corner Detection: Mathematics

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

Local quadratic approximation of  $E(u,v)$  in the neighborhood of  $(0,0)$  is given by the *second-order Taylor expansion*:

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

# Corner Detection: Mathematics



$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about  $(0, 0)$ :

$$E(u, v) \approx E(0, 0) + [u \ v] \begin{bmatrix} E_u(0, 0) \\ E_v(0, 0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0, 0) & E_{uv}(0, 0) \\ E_{uv}(0, 0) & E_{vv}(0, 0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E_u(u, v) = \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_x(x+u, y+v)$$

$$E_{uu}(u, v) = \sum_{x,y} 2w(x, y) I_x(x+u, y+v) I_x(x+u, y+v) \\ + \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xx}(x+u, y+v)$$

$$E_{uv}(u, v) = \sum_{x,y} 2w(x, y) I_y(x+u, y+v) I_x(x+u, y+v) \\ + \sum_{x,y} 2w(x, y) [I(x+u, y+v) - I(x, y)] I_{xy}(x+u, y+v)$$

# Corner Detection: Mathematics



$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about  $(0,0)$ :

$$E(u, v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0,0) = 0$$

$$E_u(0,0) = 0$$

$$E_v(0,0) = 0$$

$$E_{uu}(0,0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0,0) = \sum_{x,y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0,0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_y(x, y)$$

# Corner Detection: Mathematics



$$E(u, v) = \sum_{x,y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Second-order Taylor expansion of  $E(u, v)$  about  $(0, 0)$ :

$$E(u, v) \approx [u \ v] \begin{bmatrix} \sum_{x,y} w(x, y) I_x^2(x, y) & \sum_{x,y} w(x, y) I_x(x, y) I_y(x, y) \\ \sum_{x,y} w(x, y) I_x(x, y) I_y(x, y) & \sum_{x,y} w(x, y) I_y^2(x, y) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(0, 0) = 0$$

$$E_u(0, 0) = 0$$

$$E_v(0, 0) = 0$$

$$E_{uu}(0, 0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_x(x, y)$$

$$E_{vv}(0, 0) = \sum_{x,y} 2w(x, y) I_y(x, y) I_y(x, y)$$

$$E_{uv}(0, 0) = \sum_{x,y} 2w(x, y) I_x(x, y) I_y(x, y)$$

# Corner Detection: Mathematics



The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a *second moment matrix* computed from image derivatives:

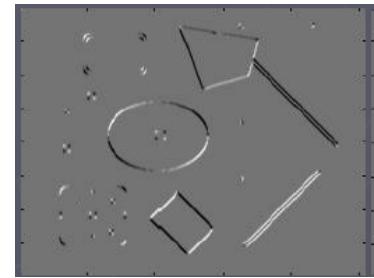
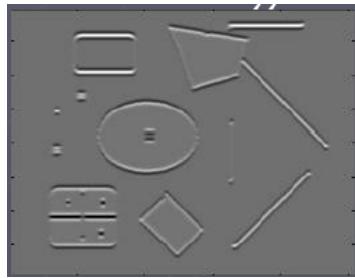
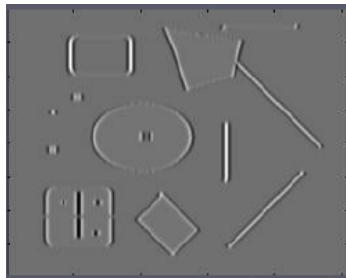
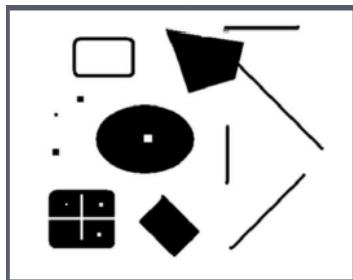
$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

# Corners as distinctive interest points

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

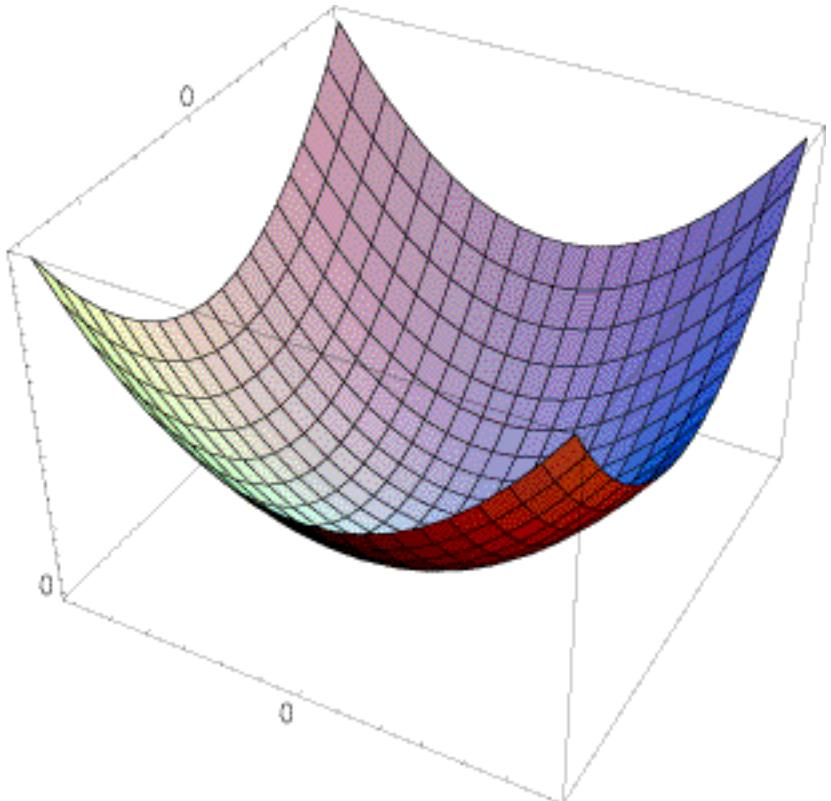
$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# Interpreting the second moment matrix

The surface  $E(u,v)$  is locally approximated by a quadratic form. Let's try to understand its shape.

$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

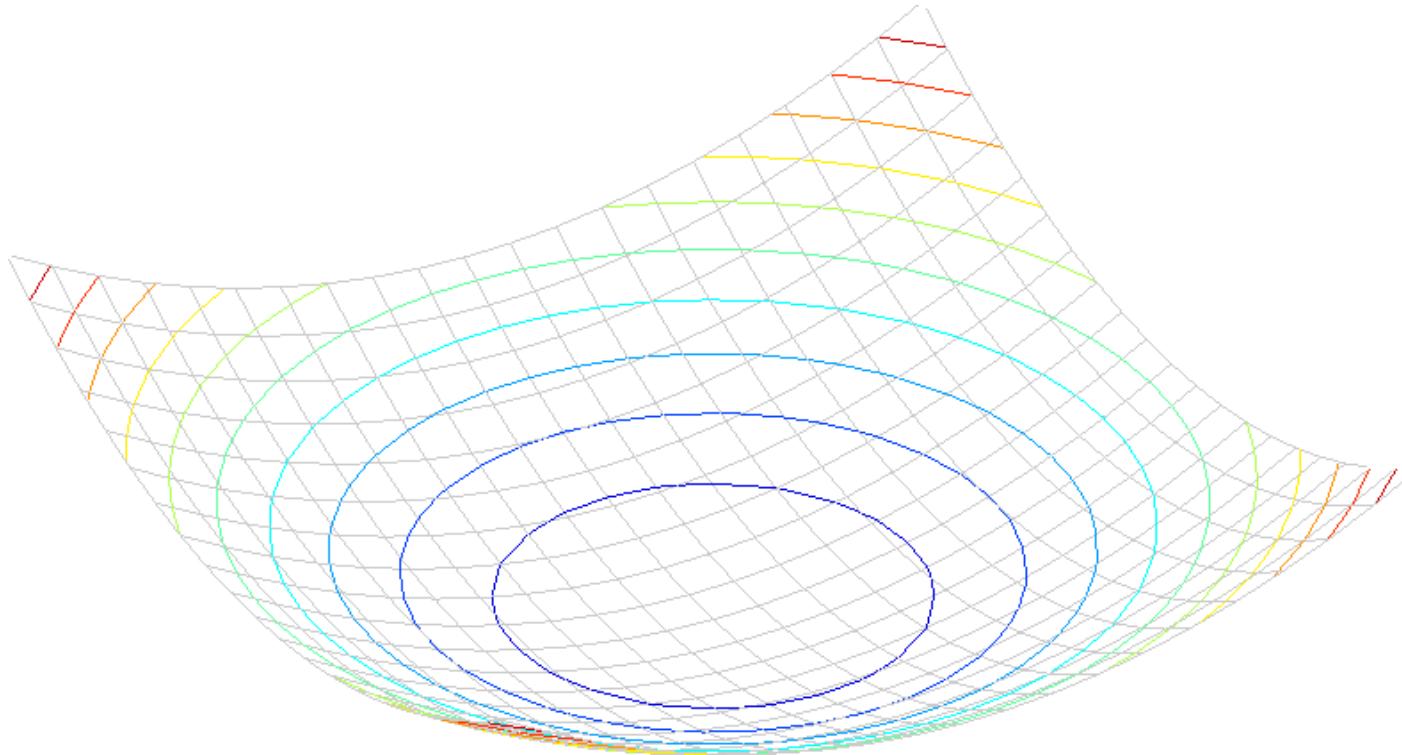




# Interpreting the second moment matrix

Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.





# Interpreting the second moment matrix

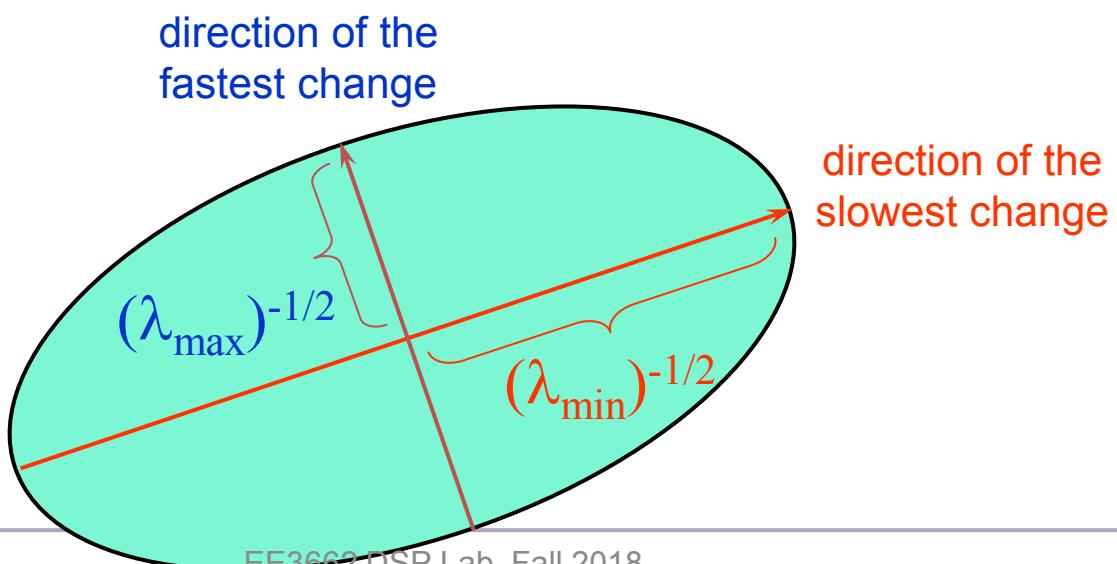
Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.

Diagonalization of  $M$ :

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$



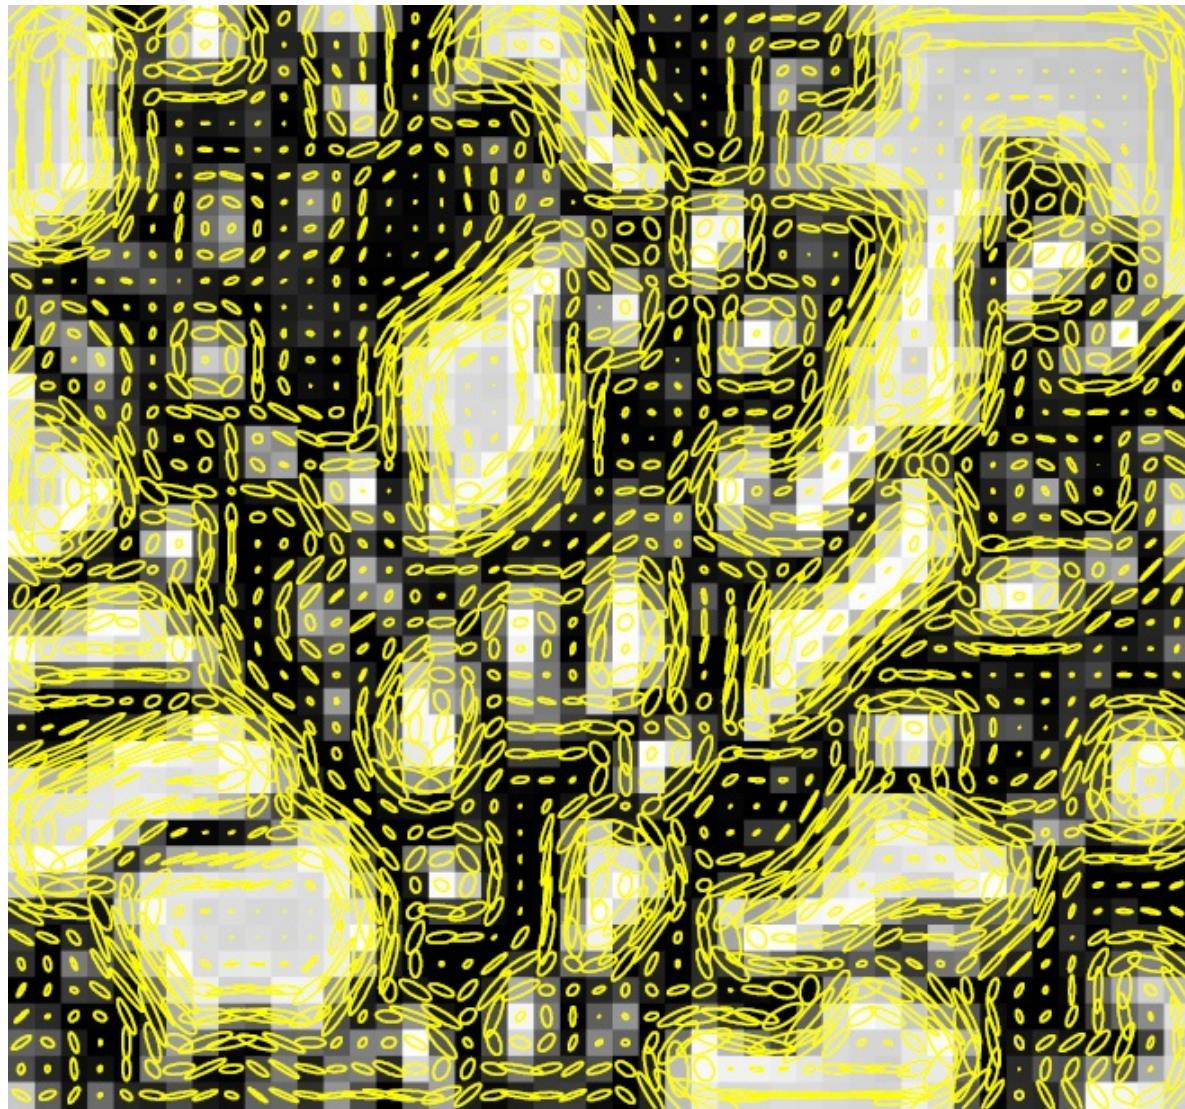


# Visualization of second moment matrices



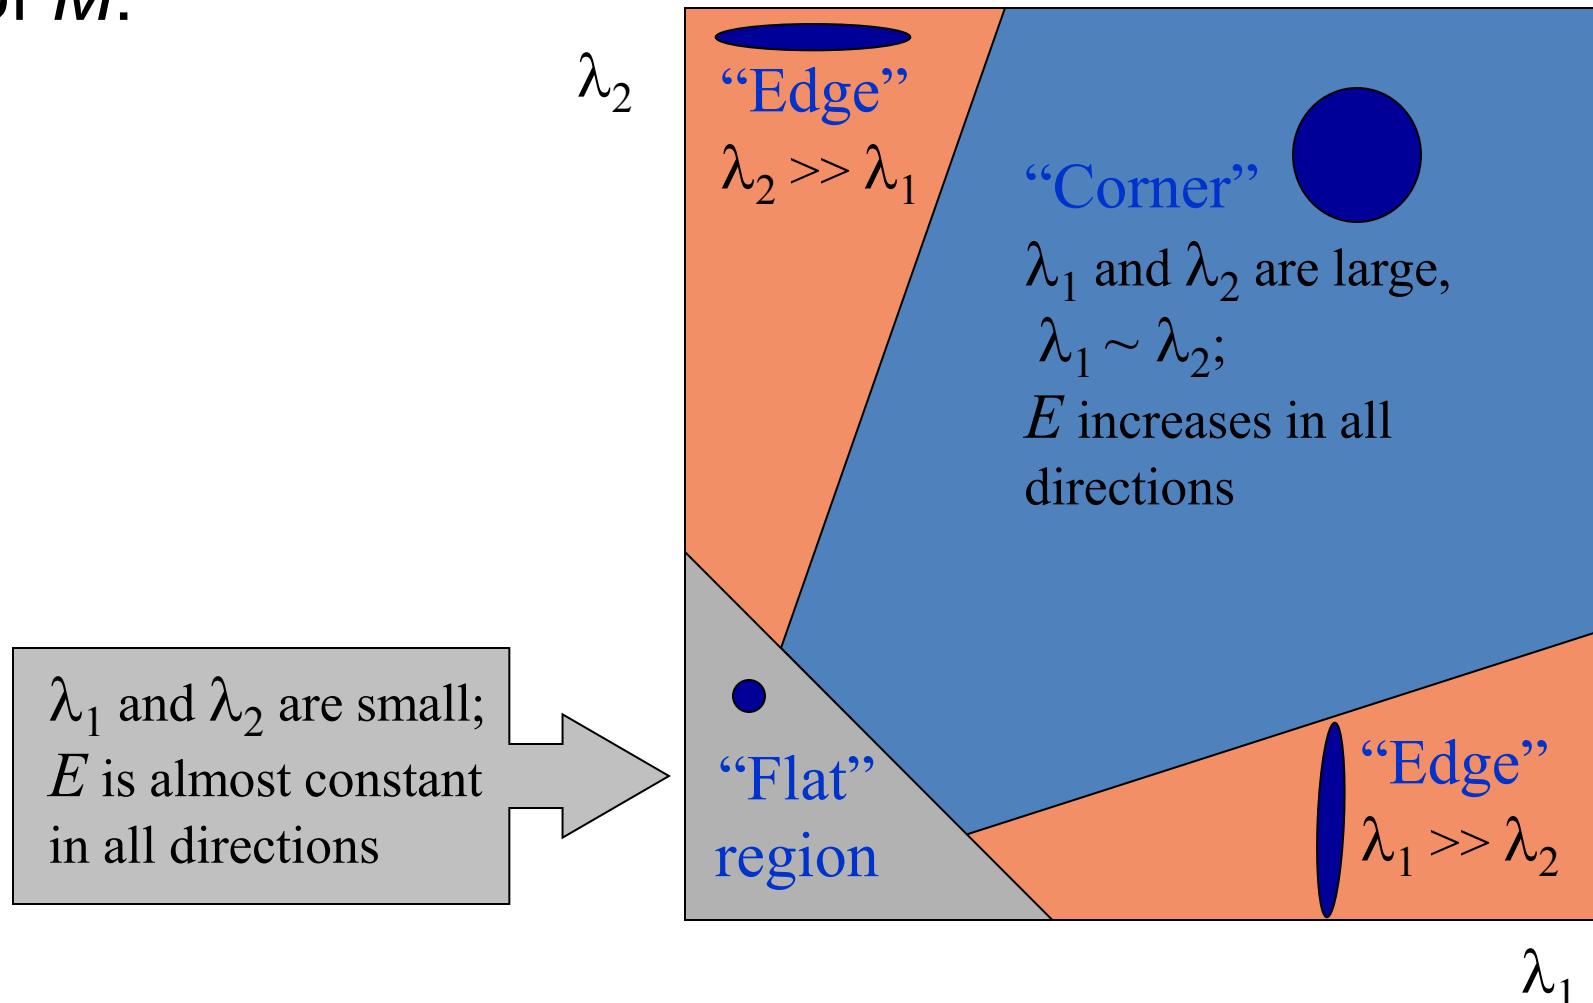


# Visualization of second moment matrices



# Interpreting the eigenvalues

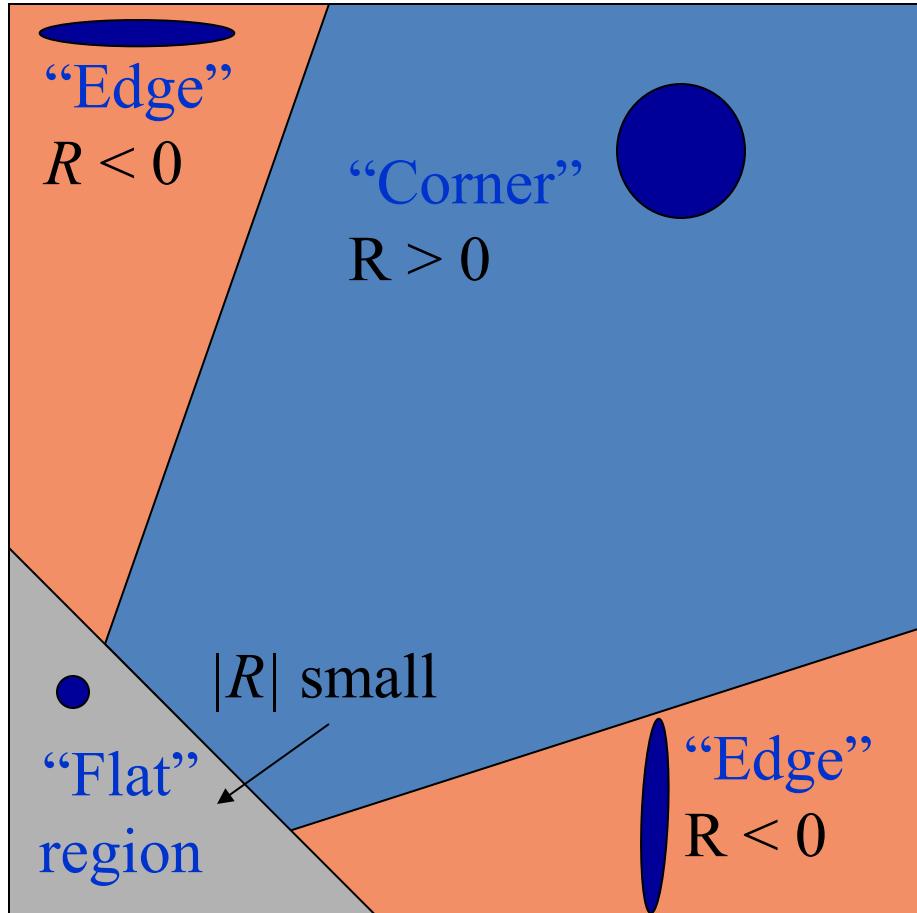
Classification of image points using eigenvalues of  $M$ :



# Corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : constant (0.04 to 0.06)





# Harris corner detector

- 1) Compute  $M$  matrix for each image window to get their *cornerness* scores.
- 2) Find points whose surrounding window gave large corner response ( $R >$  threshold)
- 3) Take the points of local maxima, i.e., perform non-maximum suppression

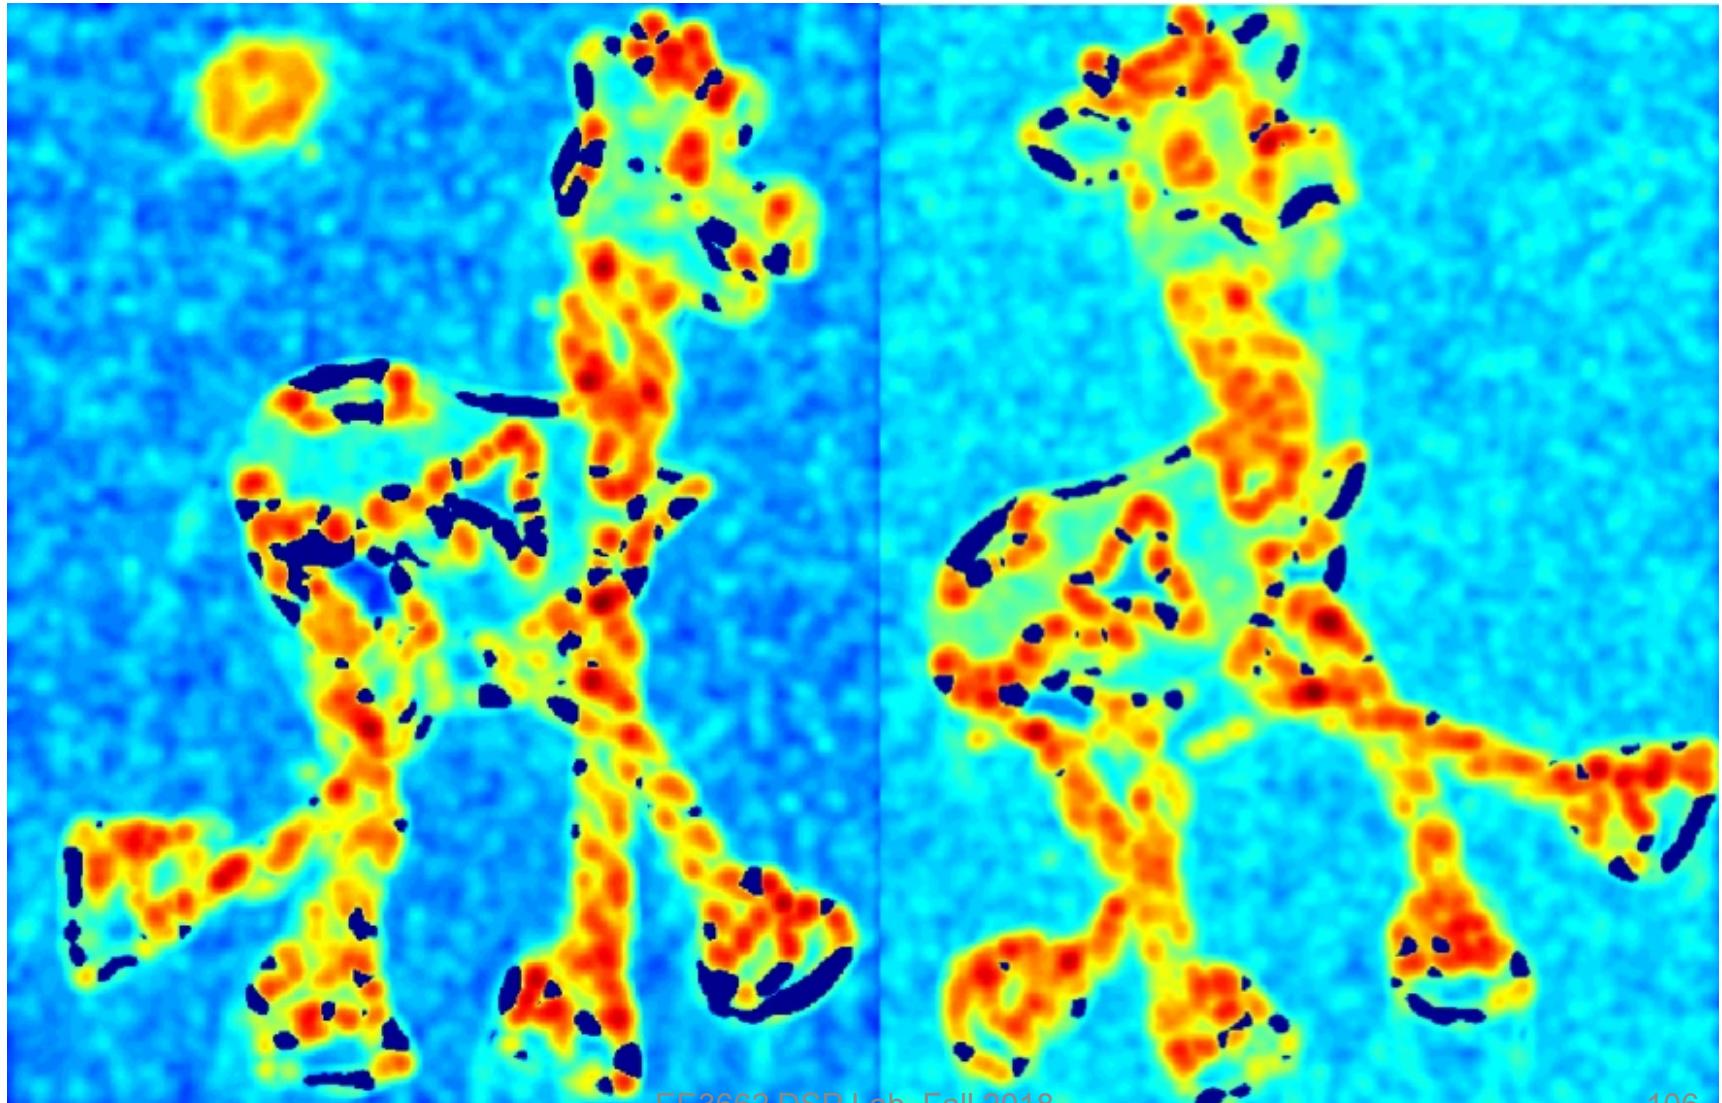
C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)" *Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

# Harris Detector: Steps



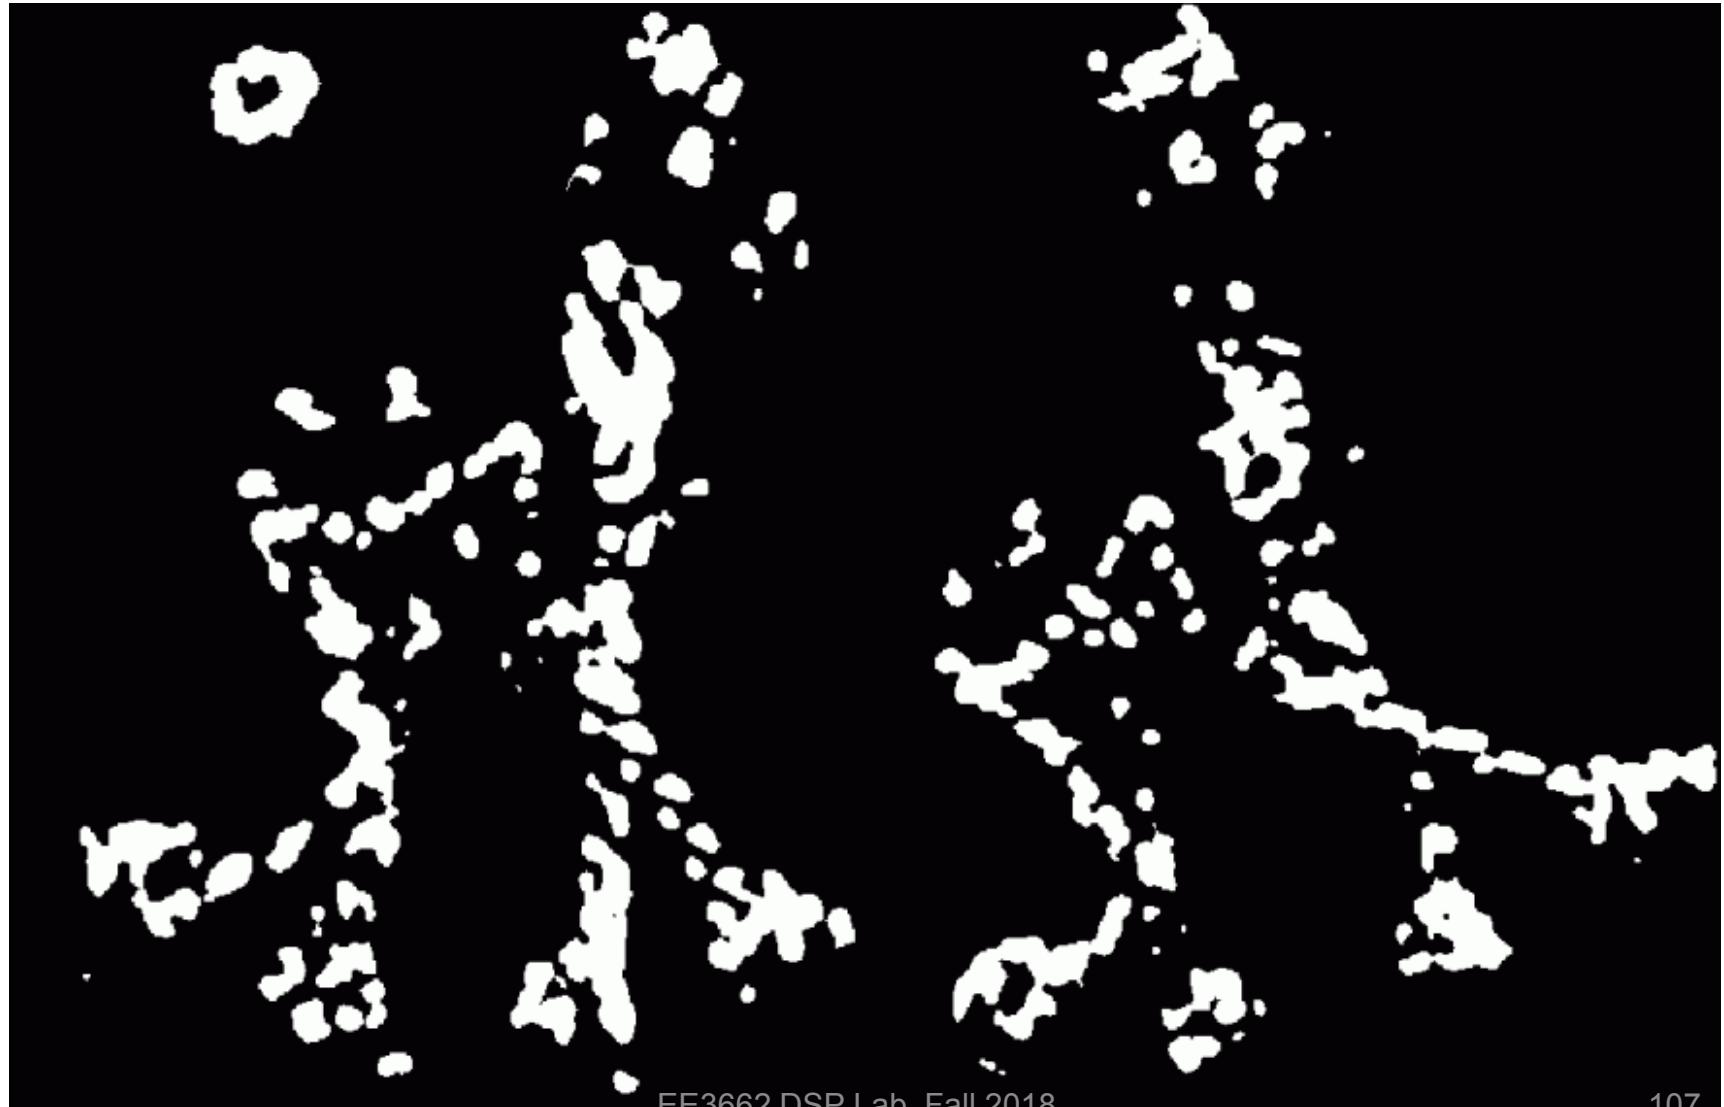
# Harris Detector: Steps

Compute corner response  $R$



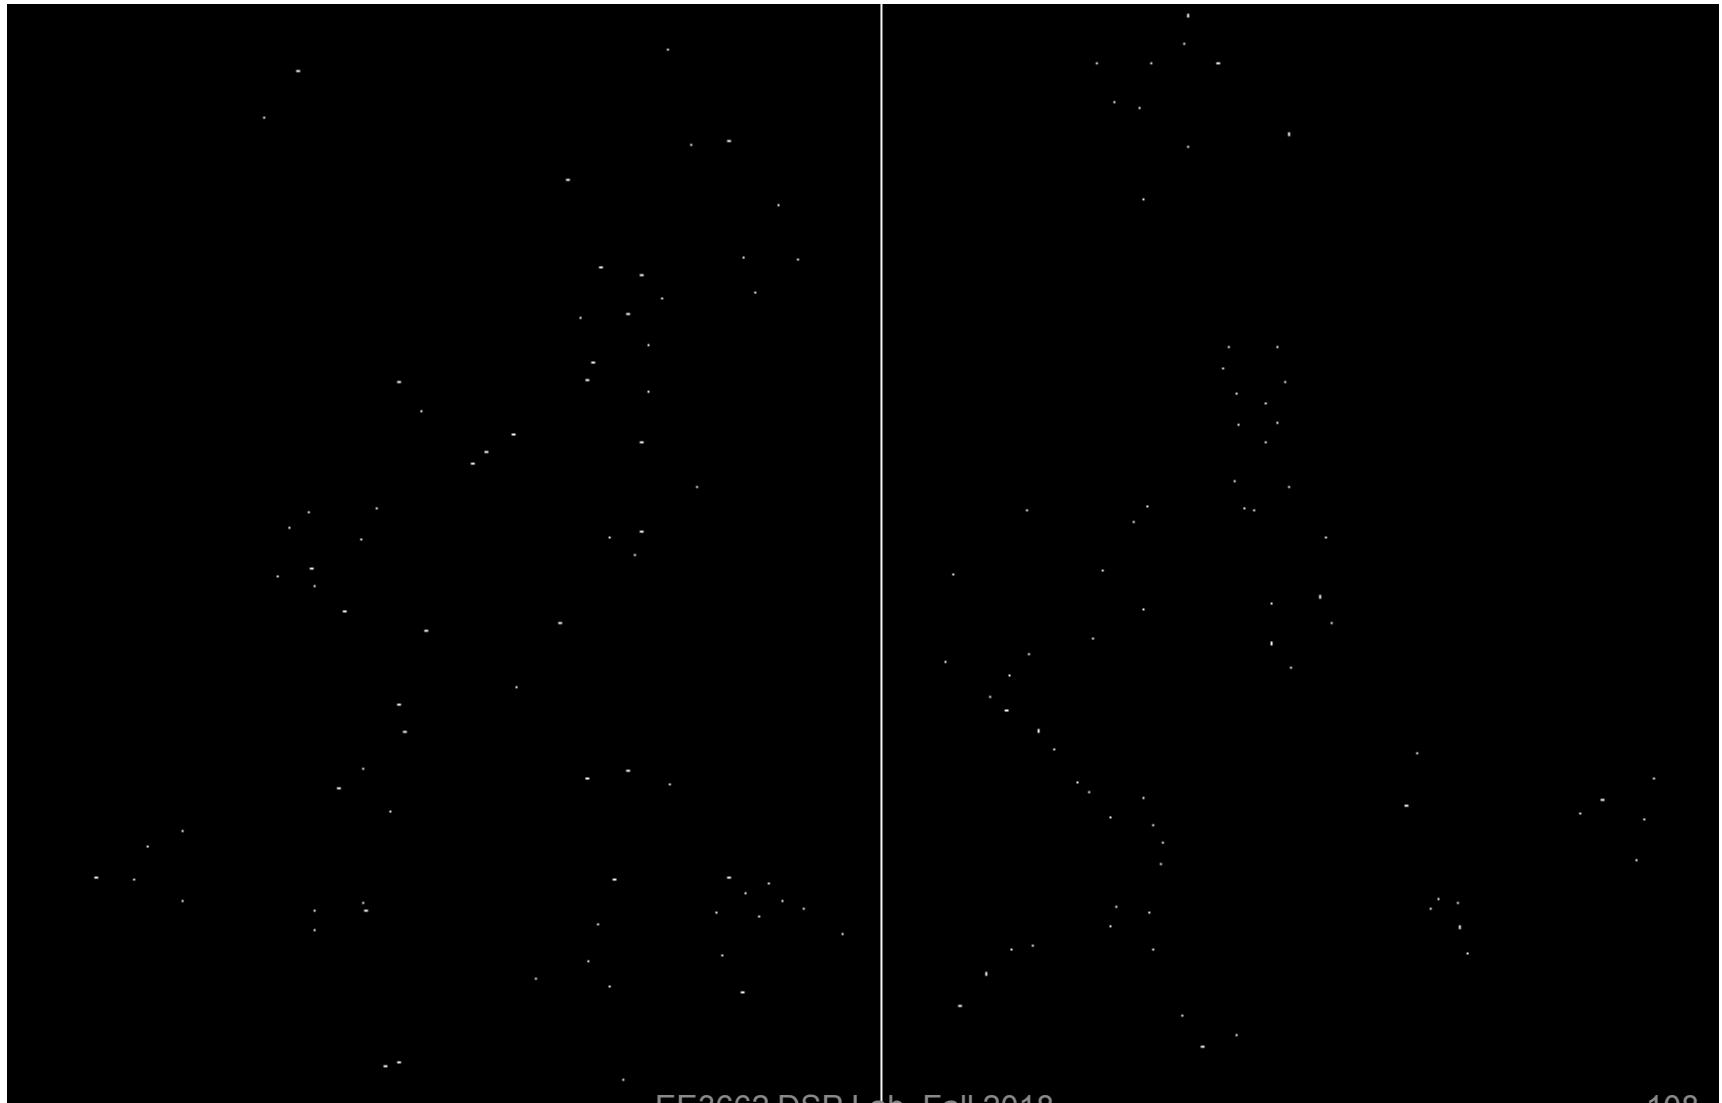
# Harris Detector: Steps

Find points with large corner response:  $R>\text{threshold}$



# Harris Detector: Steps

Take only the points of local maxima of  $R$



# Harris Detector: Steps



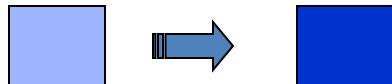


# Invariance and covariance

- We want corner locations to be *invariant* to photometric transformations and *covariant* to geometric transformations
  - **Invariance:** image is transformed and corner locations do not change
  - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations

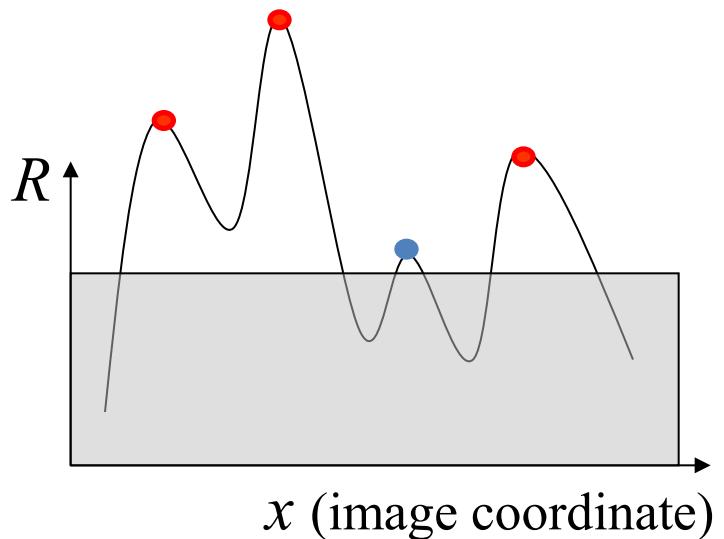
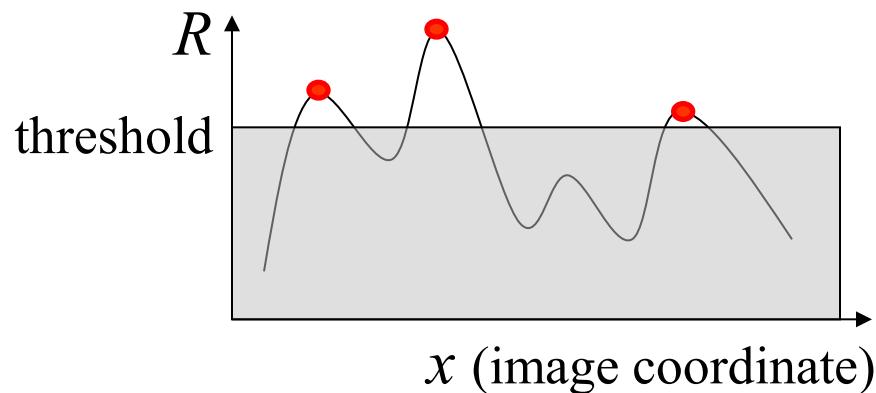


# Affine intensity change



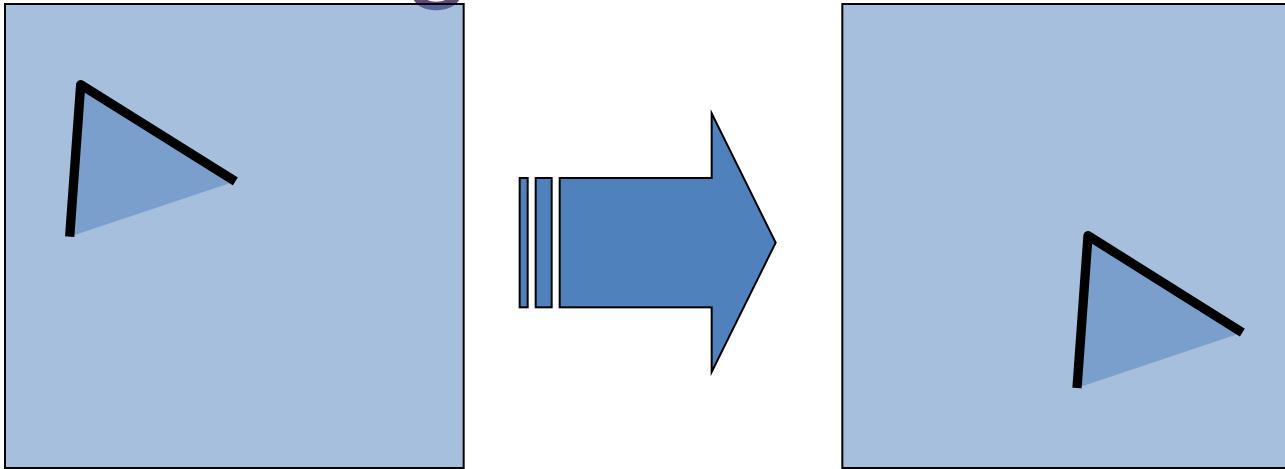
$$I \rightarrow a I + b$$

- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

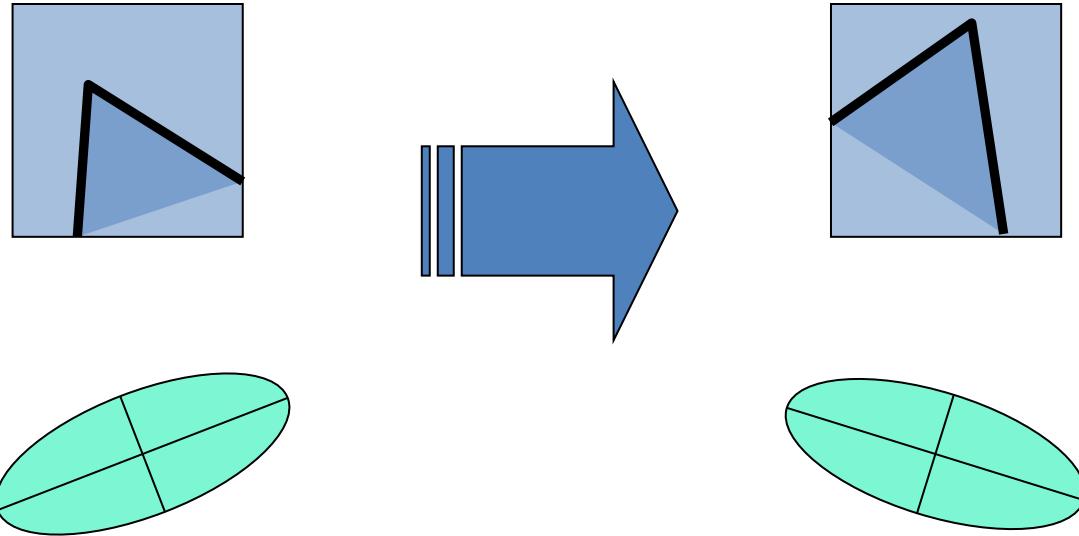
# Image translation



- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation

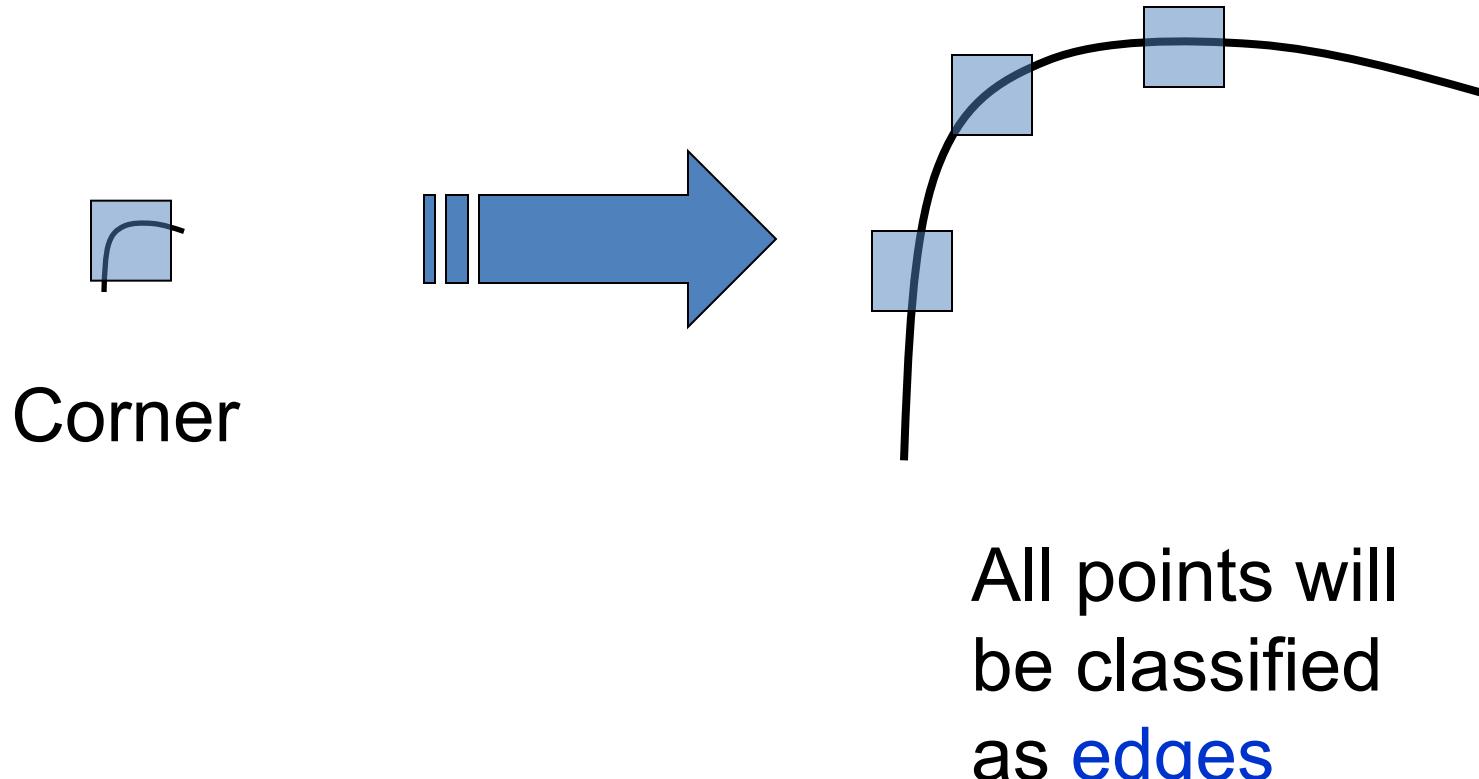
# Image rotation



Second moment ellipse rotates but its shape  
(i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

# Scaling



Corner location is not covariant to scaling!

# Lab 7: Image Filtering and Corner Detection



(Left: original image. Right: image overlap with corners.)



# Lecture Outline

- Image and Video Basics
- Recap on Image Filtering
- Interest Point
- Seam Carving

# Seam Carving for Content-Aware Image Resizing



Seam carving



Conventional scaling

S. Avidan and A. Shamir. "[Seam Carving for Content-Aware Image Resizing](#),"  
*ACM Trans. Graph.*, 2007.

# Energy function



$$e_1(\mathbf{I}) = \left| \frac{\partial}{\partial x} \mathbf{I} \right| + \left| \frac{\partial}{\partial y} \mathbf{I} \right|$$

Energy for pixel I

# Energy-based pixel removal



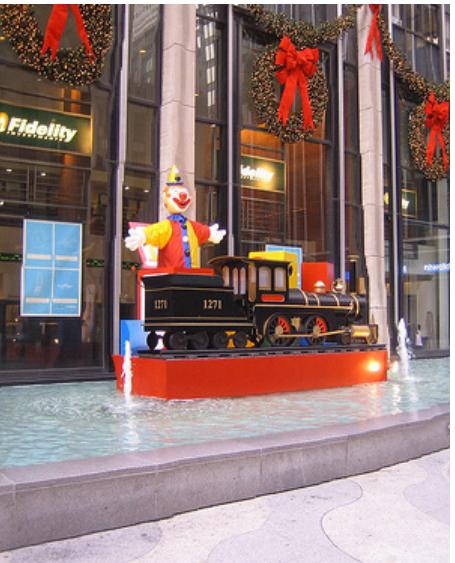
Crop



Remove the lowest-energy pixels per row



Remove the lowest-energy columns



Remove the lowest-energy seams



# Seam carving

- Define seams

$$\mathbf{s}^x = \{s_i^x\}_{i=1}^n = \{(x(i), i)\}_{i=1}^n, \text{ s.t. } \forall i, |x(i) - x(i-1)| \leq 1$$

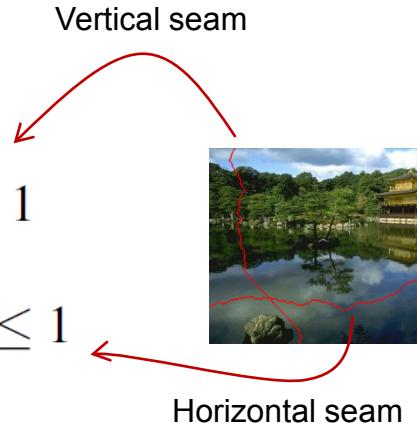
$$\mathbf{s}^y = \{s_j^y\}_{j=1}^m = \{(j, y(j))\}_{j=1}^m, \text{ s.t. } \forall j, |y(j) - y(j-1)| \leq 1$$

- Find optimal seam

$$s^* = \min_{\mathbf{s}} E(\mathbf{s}) = \min_{\mathbf{s}} \sum_{i=1}^n e(\mathbf{I}(s_i))$$

- through dynamic programming

$$M(i, j) = e(i, j) + \min(M(i-1, j-1), M(i-1, j), M(i-1, j+1))$$



# Extension

- Seam insertion



- Object removal

